



PRATIMAI • AI HANDOUTS

Build the AI Cycle

*My case file for planning, testing and improving an AI
people can trust*

CLASS 8 • LEVEL 3 — BUILD THE CYCLE

17 PAPER MISSIONS • 5 CHAPTERS • NO SCREEN NEEDED

PRATIMAI Curriculum Team

How do I know? • Who is missing? • What is the trick here?

THIS CASE FILE BELONGS TO

name • class • roll number

CASE OPENED ON



CASE CLOSED ON

License

This handbook is © 2026 the PRATIMAI Curriculum Team. Student-facing content — missions, stories, activities and illustrations-as-diagrams — is licensed under the **Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International** licence (CC BY-NC-SA 4.0): you may share and adapt it for classroom use with attribution, for non-commercial purposes, under the same licence. The Typst source files are released under the MIT licence so that any school can rebuild, translate and localise the book for its own learners.

Every activity in this book is **unplugged**: no task requires a device, an account, a photo, a voice recording or any personal information. All examples use invented people and places.

THE LICENCE IN PLAIN WORDS

You MAY	You may NOT
photocopy pages for your class or study group	sell this book or any copy of it
adapt missions for your learners and share them	remove the credit lines or the licence
translate the book and keep the same licence	use it to advertise any product or service
point other schools to the free source files	claim the activities are your own invention

Series: PRATIMAI · AI Handouts
Title: Build the AI Cycle — Class 8 · Level 3 · BUILD THE CYCLE
Author: PRATIMAI Curriculum Team
Publish Date: 2026-10-10
Published by: PRATIMAI | AI-Schools
Edition: Machiatto Edition 1.0

Set in **Nunito** and **Baloo 2** on warm paper stock, typeset with Typst using the **Machiatto** template — which implements the MoKa Reads publication specification: title page, license, acknowledgements, preface, contents, then chapters opening with a summary and a mini table of contents. Both typefaces are used under the SIL Open Font Licence; the template ships under the MIT licence. Every chart, grid and pixel drawing in this book is generated in pure Typst code — no images were imported, so every page prints crisply at any size.



PRATIMAI • AI HANDOUTS

Acknowledgments

To the Classes 6 and 7 investigators whose case files came back covered in pencil, arrows and honest confidence circles — this book is built on your evidence. Your teachers told us which missions turned a quiet classroom into a planning meeting, where the instructions needed one more worked example, and which questions made the whole class argue productively for ten minutes. Every note shaped a page here.

To the curriculum reviewers who tested every mission for the “unplugged” promise: thank you. You caught the draft that quietly wanted a laptop for the proposal task, and you were right. An AI builder at this level needs paper, pencils, sticky notes, playing cards and teammates — the same tools professional teams use on their first planning day. And to the grown-ups who will be asked “who is responsible when the machine gets it wrong?” over dinner: this book trains the child to answer you with reasons, not slogans.

Thanks as well to the open-source community behind Typst and the Machiatto template that gives this book its typesetting, and to the designers of Nunito and Baloo 2, whose friendly letterforms keep pages of hard planning feeling light. Finally, to the project lead holding this book: every blank table, empty checklist and half-drawn cycle map in here was left open on purpose — because on a real project, the plan is never handed to you. The next plan in this file is yours.

17 PAPER MISSIONS • 16 WORDS TO KEEP • 5 CHAPTERS • 1 BADGE EARNED

How to Use This Handbook

Welcome back, Detective! Two years of clues sit behind you: in Class 6 you learned to **spot** learning machines, and in Class 7 you operated one — sorting, predicting and watching machines get fooled. This year you take the next step: you **plan and build**. Every serious AI system in the world began as a plan on paper — a sharp problem statement, a list of data, a test, a fair check — exactly the artefacts you will produce in these five chapters. You will still not need a computer, a phone or an app: professional builders plan on whiteboards and paper before they touch a keyboard, and so will you. Every chapter hands you missions (**tasks**), evidence to examine, and empty spaces to fill with your decisions. This book belongs to you: write, draw, cross out, improve.

1 READ

Read the mission twice.
Underline anything that puzzles you — puzzles are clues.

2 DO

Do the task with your partner, group, class or family, exactly as the labels say.

3 WRITE

Write what you decided in the spaces, then shade your confidence circles honestly.

The work-mode labels

ALONE	your own case — you investigate by yourself
PAIR	two detectives — share one book, take turns reading codes aloud
GROUP	three to four detectives — decide roles before you start
CLASS	the whole class plays together, led by your teacher
HOME	a mission for home — show a grown-up what you discovered

HOW TO READ A MISSION

T8-01 Mission title

PAIR • 15 MIN

■ **QUICK WIN** sits at the right edge of the easier missions.

Every mission has a code like **T8-01** — **T8** means this Class 8 handout, and **01** is the mission number, so your teacher can say “open Task T8-09”. The **PAIR** label tells you who to work with, **15 MIN** is your time limit, and **QUICK WIN** marks an easy first success. This book’s missions ask you to **decide and justify**, not just to sort — so expect to write a reason beside almost every answer. Nobody fails in this book, because every attempt teaches a clue.

THE QUESTION HABIT — YOUR SUPERPOWER

Your three detective questions — “**How do I know?**”, “**What is missing?**”, “**Who made this?**” — get two powerful companions this year. For every plan you meet (including your own): “**Who is missing?**” — which group was never asked, never counted, never helped? And for every claim a machine or a company makes: “**What is the trick here?**” Builders who ask these two questions early catch unfair machines **before** people get hurt. Use all five questions in this book, at home, and on every app that promises to be smart.

A Note for Grown-Ups

This handout builds **AI understanding** — how machines learn from data and how humans should judge them — through 17 paper-and-pencil missions. It is fully **unplugged**: no task requires a device, an account, a photo, a voice recording or any personal information. This year the big ideas are the AI project cycle (scope, data, build and test, reflect and improve), data representation and sampling, consent, a nearest-example classifier run by hand, accuracy, no-code tool types and their limits, how text generators predict next words and why they invent, and responsible-use decisions. The best help you can give is to ask the Question-Habit questions, praise careful **reasoning** more than quick right answers, and treat the confidence circles as honest thinking, not scores. A separate teacher pack carries answer notes and rubrics; this book deliberately never prints solutions beside a task — only short hints at the very back.

Four ways to help — without giving answers

1 Ask about the plan

When your builder is stuck, ask: **“Who exactly has this problem, and what data would show it?”** A sharp question rebuilds a plan faster than a hint.

2 Praise the trade-off thinking

Swap “correct!” for **“you noticed the plan helps one group and misses another”** or **“you chose the slower option for a fair reason — that is engineering.”**

3 Let the plan fail

A plan that fails its test is the most valuable page in the book. Ask **“what will you change in version 2?”** — that loop is the whole subject.

4 Connect to real decisions

Point at systems around you — the bus route app, the exam result predictor, the clinic’s queue machine — and ask **who is missing** from each one’s data.

IF YOUR CHILD ASKS...


“Can I build real AI after this book?” — Say: you will have built the **plan** the way professionals start — problem, data, test, fairness check. The coding comes years later, and it is the easy-looking part. **“Why does the chatbot make things up?”** — Say: it writes by guessing likely next words, not by checking facts, so a confident lie and a confident truth are built the same way. That is why this book trains verify-first habits. **“Is no-code cheating?”** — Say: no. Choosing the right tool, the right data and the right test is the thinking; the tool is just the last step.

SIGNS OF BUILDER GROWTH WORTH PRAISING

Watch for these quiet changes — they matter more than any correct answer: your child **states the problem precisely** before jumping to solutions · they ask **whose data** is in a system and **whose is missing** · they test a plan before trusting it — and improve it after · they say **it depends** and then explain what it depends on · they name a trade-off — speed against fairness, cheap against accurate — and defend a side. When you spot one, name it: “That is exactly what builders do.”

My AI Journey Map

THE FIVE STRANDS YOU WILL MEET AT EVERY STOP



THINKING-TOOLS THREAD

runs through every stop: **decompose** the problem · **test before trusting** · **improve in loops** · **optimise under limits**

THE FIVE STOPS ON YOUR CASE



BY THE END OF THIS BOOK, I CAN...

- ☐ define **AI** and **machine learning** in my own words and place chatbots, image classifiers and recommenders on a who-learns-what map
- ☐ walk a local problem through the full AI project cycle — on paper
- ☐ hand-run a nearest-example classifier, compute its accuracy on a test set, and suggest one improvement
- ☐ audit a small dataset for missing groups and compute percentages per group
- ☐ explain why collecting data needs **consent** and privacy care
- ☐ match no-code tool types — classifier, chatbot, predictor — to problems, with inputs, outputs and limits
- ☐ explain why a text generator can sound right and still be wrong — and verify its claims
- ☐ apply a fairness checklist and propose fixes for an unfair dataset
- ☐ make and defend a responsible-use decision about privacy, misinformation or social impact
- ☐ break a multi-variable problem into parts and find the best choice under limits

Contents

1. The AI Project Cycle	8
1.1 From user to builder: who learns what	9
1.2 The loop every AI project runs	12
1.3 The Plant Doctor case: a complete walk-through	14
2. Data & Fairness	18
2.1 The audit: who is inside the table?	19
2.2 The sampling game: small groups, big surprises	21
2.3 Fairness Fix-It: repairing a broken guest list	23
2.4 Consent: the door every example walks through	26
3. How Models Learn & What No-Code Tools Do	31
3.1 The nearest-example machine	32
3.2 The honest score: accuracy on a hidden test	34
3.3 The Paper Classifier Lab: build, test, fail, improve	37
4. Ethics & Responsible AI	42
4.1 The next-word machine: how text generators actually work	43
4.2 Hallucination Hunt: verify before you trust	44
4.3 The Responsible-Use Court: decisions with consequences	46
4.4 My AI-Use Pledge	47
5. Capstone: The AI Project Proposal	51
5.1 The proposal: your plan on paper	52
5.2 The Peer Review Gallery	54
5.3 Optimise It: the best you can afford	56
My AI words	60
My progress	62

The boxes in this book

You will meet these boxes on every case. Each one wants something different from you — collect them all.

MISSION

A paper mission with a code like **T6-01**. The label tells you who you work with and how many minutes you get. Do it — then write. Writing **is** the mission.

WORD POWER

A word worth keeping in your detective kit. Say it three times today — words are the tools of thinking.

SELF-CHECK

Skills to grade yourself on. Shade a confidence circle honestly — honest circles make you learn faster.

THINK & INK

A question with lines. There is no single right answer here — your **reasons** are the answer.

Also on patrol: **MYTH BUSTER** busts a wrong idea many people believe · **DETECTIVE'S NOTE** hands you a professional's secret · **HOME LINK** (dashed) is a mission to take home and teach a grown-up.

1. The AI Project Cycle

GUIDING QUESTION

What does it take to build an AI that people can trust?

You are promoted: from machine-**operator** (Class 7) to machine-**builder**. But no real AI begins with a machine — it begins with a plan. In this chapter you learn the loop every serious AI team runs: **scope** a sharp problem, **collect** the right data, **build and test** a model, then **reflect and improve**. You will turn vague complaints into precise problem statements, sequence the cycle from cards, and walk a complete Plant Doctor project end to end to find exactly where it went wrong.

MISSIONS **T8-01 – T8-03**
OUTCOMES **8.U1 · 8.L1** STRANDS

L T

Links: Maths — percentages & data · English — precise definitions · Science — plant health case

In this chapter

- 1.1 From user to builder: who learns what 9
- 1.2 The loop every AI project runs 12
- 1.3 The Plant Doctor case: a complete walk-through 14

WORDS YOU'LL MEET HERE

machine learning

AI project cycle

problem statement

dataset

WARM-UP · BEFORE YOU READ

ALONE · 2 MIN

Your school wants “an AI that helps students eat better”. Write your honest first reaction: what would you need to know BEFORE anyone builds it? Keep your answer — you will rewrite it after Task T8-01.

DETECTIVE'S TOOLKIT — YOU WILL NEED: pencil · scissors & envelope for cards · sticky notes · your Class 7 notebook

1.1 From user to builder: who learns what

Last year you trained machines with your own hands, so you already know the secret: a **learning** machine is not given rules — it is given **examples**. This year we name the big idea properly.

Artificial intelligence is the wide field: any machine that does things we used to need human judgement for. **Machine learning** is the engine inside most modern AI: the machine improves at a task by finding patterns in **examples** (a **dataset**) instead of following rules a person typed. And every AI product you meet is really a machine that learned **one specific thing** from **one specific kind of data**. A chatbot learned from enormous amounts of human text. An image classifier learned from photos with labels attached. A recommender learned from what people like you watched, bought or liked. Different food, different learning.

AT A GLANCE · THE WHO-LEARNS-WHAT MAP

The AI product	What it learned from (its data)	What it can do	What it cannot do
chatbot	huge amounts of human writing	continue a conversation, explain, draft	know whether what it wrote is TRUE
image classifier	thousands of photos, each with a label	sort new photos into its labels	explain why , or see what no photo showed
recommender	records what similar people chose	suggest the next video or product	understand you, or ask you a question

Keep this map in mind all year. Every AI system is somewhere on it — and the last column is where builders earn their salary.

TWO MACHINES, HONESTLY COMPARED

Pick one from my morning list	What it learned from	Its honest cannot-do

The who-learns-what map in ONE sentence of mine:

DETECTIVE DRILL · SORT THE SIX — WHO LEARNS WHAT?

Place each product on the map: what did it learn from, and what is its honest cannot-do? Write C (chatbot), I (image classifier) or R (recommender) — then one cannot-do of your own:

The product	C / I / R?	Its honest cannot-do
a video app that queues up cricket highlights for you		
a photo app that names the dog breed		
an exam-doubt solver that writes full explanations		
a shop bell that offers coupons based on your old bills		
a scanner that spots cracked railway panels		
an email filter that learns which mail you mark as junk		

Two of these learned from the SAME kind of data — can you spot which pair? That pair is your proof that data, not the product name, decides what a machine can do.

QUICK POLL — THE FAMILY QUESTION

Of the six products above, which ONE would your family actually pay for — and which one would they refuse on principle? Two lines, one reason each:

PREDICT THE LEARNING Which ONE product above do you predict learned from the FEWEST examples — and what could go wrong for it? Two lines:

T8-01 Problem Statement Sharpener

■ QUICK WIN PAIR • 15 MIN

Real projects die from **vague problems**. “Our village has a water problem” cannot be built; “detect which public taps run dry for more than two days, so the ward office can dispatch repairs” can. The sharpening tool is the **4W test** — **Who** exactly has the problem? **What** exactly goes wrong? **Where** does it happen? **Why** does it matter? Take these three blunt complaints and sharpen each one into a problem statement a builder could actually plan around:

The blunt complaint	Your sharpened problem statement (use all four Ws)
“Students lose their notebooks.”	
“The bus is always late.”	
“People waste water.”	

Now the hardest step: cross out one W in your bus statement and watch it wobble. Which W was load-bearing — remove it and the plan collapses?

WORD POWER • 1

Definition: **problem statement**

One precise sentence saying exactly what problem you solve, for whom, where.

Say it three times today — it sticks!

MY MORNING MACHINES — name three “smart” helpers you met before school:

1. _____ 2. _____ 3. _____


By page 20 you will be able to say which of the three learned from examples — and what its data was. Watch your list change.

DETECTIVE'S NOTE

A good problem statement sounds boring — and that is a feature. “Reduce” and “improve” are dreams; “detect”, “count”, “predict X for group Y in place Z” are plans. In the capstone of this book you will write one of these for a real problem you choose, and your classmates will pressure-test it. Sharpen now; shine later.

1.2 The loop every AI project runs

Here is the whole subject of this book in four boxes. First, **problem scoping**: turn a vague worry into a sharp statement (you just did this). Second, **data**: find or collect examples the machine can learn from — and check who is inside them. Third, **model building and testing**: train on some examples, test on hidden ones, and measure honestly. Fourth — the step beginners skip — **reflection and improvement**: the first version is always wrong somewhere; the team reads the failures, fixes the data or the design, and loops back. Notice the arrows: this is a **cycle**, not a straight road. Professional teams go around it many times before anything ships.



THE LOOP IN ONE MINUTE Explain the diagram to a Class 6 student in exactly two sentences — one for the boxes, one for the arrows:

T8-02 Cycle Cards

GROUP • 15 MIN

Your teacher has an envelope with five cards: **PROBLEM, DATA, MODEL, TEST, IMPROVE**. (No envelope? Copy the five cards onto paper slips.) As a team, lay them out in the order a real project runs, then defend your order — one line per card:

Card	Position we chose (1st, 2nd...)	Why it sits there — one line
PROBLEM		
DATA		
MODEL		
TEST		
IMPROVE		

The twist question: after IMPROVE, which card does the team return to? Draw the loop arrow on your table — and explain in one line why the cycle never really “ends” for a machine that keeps meeting new people:

WORD POWER • 2

Definition: **AI project cycle**

The loop: scope the problem, collect data, build and test, reflect and improve.

Say it three times today — it sticks!

WORD POWER • 3

Definition: **dataset**

The collection of examples a machine learns from, or is tested on.

Say it three times today — it sticks!

DETECTIVE DRILL · NAME THE STAGE

Each sentence below comes from some AI project. Write which stage it belongs to: P (problem), D (data), M (model), T (test) or I (improve):

What the team did or said	Stage?
"We spent two days deciding: are we predicting late buses or crowded buses ? Pick one."	
"We asked 60 commuters to log their wait times for two weeks."	
"We taught the machine on 40 logged trips, keeping 20 trips hidden."	
"On the hidden trips it was right 13 times out of 20 — let's look at the 7 misses."	
"The misses were all rainy-day trips. Version 2 gets a weather column."	
"Wait — before any of this: WHO exactly suffers from late buses, and where?"	

■ MYTH BUSTER

MYTH "AI = ChatGPT."

TRUTH A chatbot is one famous resident of the AI city, not the city itself. The spam filter in your mail, the crop-disease scanner in a field office, the queue predictor at a clinic — all AI, none of them chat. And most of them learned their skill from a carefully chosen **dataset** through the very cycle you just sequenced. When someone says "AI", train yourself to ask: **which machine, learned from what data, tested how?** That question is the whole of this book.

1.3 The Plant Doctor case: a complete walk-through

Time to watch the whole cycle run — and fail. Read this true-to-life case from the fictional village of Amravati. A team built an AI to diagnose plant disease from leaf photos. Every stage was done with enthusiasm. Your mission is to find where it **actually** went wrong, because the report card said only: "accuracy 82%."

PREDICT THE FAILURE — BEFORE YOU READ THE DIARY

An AI project done "with enthusiasm" by one person in one week usually fails at the same stage. Which stage do YOU predict, and what is your evidence from the sentence above?

Read the project diary, then answer as a team.

PROJECT DIARY — PLANT DOCTOR v1

Stage 1 — Problem: “Help farmers know if a plant is sick.” (The team liked this slogan so much they kept it.) **Stage 2 — Data:** A member collected 2,000 leaf photos in one week — all from **his own farm’s hybrid tomato field**, taken at **noon in full sun**, all labelled by **one nephew** who “knows plants”. **Stage 3 — Model & Test:** They trained on 1,600 photos and tested on 400 hidden ones. Accuracy: 82%. **Stage 4 — Deploy:** The app was announced at the market. Two weeks later the complaints arrived: it called every diseased leaf “healthy” on the small farms near the river, and it knew nothing about the local desi (native) varieties that look completely different from hybrid tomatoes.

Diary stage	One thing that was done too fast or too narrowly	What the team should have done instead
1 · Problem		
2 · Data		
3 · Model & test		
4 · Reflection		

The accuracy trap: 82% sounds like a good score. Write two reasons why that single number hid the failure completely:

Whose problem statement would you write instead? One sharpened sentence for the REAL users:

Self-check

shade one: first circle = not yet · middle = almost · last = got it

- ☐ I can turn a vague complaint into a problem statement with Who, What, Where and Why ○○○
- ☐ I can name the four stages of the AI project cycle and explain why it loops ○○○
- ☐ I can place a chatbot, an image classifier and a recommender on the who-learns-what map ○○○
- ☐ I can point at the Plant Doctor diary and name the stage that failed first ○○○

■ THINK & INK

A “smart” app or service my family uses was really built for someone else. Its problem statement was probably written for ..., because ...

DETECTIVE DRILL · SHARPEN OR SPOIL?

Two problem statements enter the workshop. Mark each S (sharp enough to plan with) or V (still vague), then fix the vague one:

Problem statement	S or V?	If V — the missing W
“Predict which water pumps in Ward 12 will fail this monsoon, so repairs happen before the rains.”		
“Make education better with AI.”		
“Flag Class 8 homework sheets that arrive blank, so the teacher can check 3 instead of 40.”		
“Help people.”		

CASE NOTES — WHAT CHANGED IN MY THINKING?

Before this chapter I believed building AI starts with the machine. Now I know it starts with ...

WHERE WOULD YOU ENTER THE LOOP? A friend's school already has "students forget things". Which cycle stage would YOU enter the loop at — and what would you build your version on?

CHAPTER 1 CLUES — POCKET THESE

- AI is the field; **machine learning** is the engine that learns patterns from a **dataset**.
- A **problem statement** passes the 4W test: Who, What, Where, Why — all four, or it wobbles.
- The project cycle loops: **problem → data → model & test → reflect & improve → back again**.
- One accuracy number can hide **who** the machine fails — the Plant Doctor proved it.

The clue I would tell my family tonight:

Chapter 1 Checkpoint one line each — no peeking back!

ALONE • 3 MIN

1. A complaint says "fix the school toilets." Write its sharpest one-line upgrade using at least three of the four Ws.

2. The Plant Doctor team tested on 400 hidden photos and scored 82%. Name the TWO diary stages that deserved the blame anyway.

3. A chatbot and an image classifier — which one learned from human writing, and what is the other one's honest cannot-do?

MY CASE JOURNAL — TODAY'S SHARPEST CLUE ABOUT THE PROJECT CYCLE

2. Data & Fairness

GUIDING QUESTION

If the data only knows some of us, who does the machine actually serve?

A machine is a mirror of its dataset — show it only some people and it serves only some people. In this chapter you audit real tables like a data detective (counting who is over-represented and who is missing), run a sampling game that ends with a surprise, repair an unfair dataset three different ways, and write the consent rules every honest collector must follow. This is the chapter where builders learn that “the data said so” is never the end of the story.

MISSIONS T8-04 – T8-07
OUTCOMES 8.D1 • 8.D2 • 8.R1 **STRANDS**

D R

Links: Maths — percentages, sampling · Social Science — inclusion & community · English — persuasive asking

In this chapter

2.1 The audit: who is inside the table?	19
2.2 The sampling game: small groups, big surprises	21
2.3 Fairness Fix-It: repairing a broken guest list	23
2.4 Consent: the door every example walks through	26

WORDS YOU'LL MEET HERE

sampling consent fairness representation

WARM-UP • BEFORE YOU READ

ALONE • 2 MIN

If a “smart attendance” camera was trained only on morning-shift students, who would it fail to recognise — and why would that matter? Write one line.

DETECTIVE'S TOOLKIT — YOU WILL NEED: pencil · calculator or rough column · class register (or a made-up list of 24 names) · sticky notes

2.1 The audit: who is inside the table?

Every dataset is a guest list. Somebody decided who got invited — which streets, which ages, which languages, which crops — and everybody left off the list simply does not exist **for the machine**. Data scientists call the balance of that guest list **representation**: does the dataset look like the real group of people the system will meet? You already know what happens when examples are few or one-sided (Class 7, T7-05). This year we go further: we **count**. A fair audit does not say “we should be more inclusive” — it says “40 of our 50 examples are from one ward; that ward is 20% of the town; we are overweighted by a factor of four.” Numbers turn a feeling into a plan.

GUEST LIST OF MY CLASSROOM — A 60-SECOND PRACTICE AUDIT

If a machine were trained only on the students sitting in your row or group right now, which TWO voices or habits from the wider class would it never learn?

QUICK COMPUTE — WARM UP THE AUDIT MUSCLE

If the dataset has...	...and this many are from Ward 3	What percentage is that?
40 examples	10	
200 records	120	
50 photos	2	

Same three numbers, three different weights — percentages are how an audit speaks out loud.

TWO-MINUTE VOLLEY — SAY IT OUT LOUD WITH A PARTNER

Take turns, one sentence each: “A dataset is a guest list because ...” — then “Consent matters because ...” — then “A fix without a price is ...” Two lines for the best sentence you heard:

A city clinic built a “fever triage” helper trained on 200 patient records. Before you audit: **predict first**. Which group do you EXPECT to be overweighted in a clinic’s own records — and why? Write one line, then audit for real:

THE 200 TRAINING RECORDS

Group in the records	Records
Adults, Ward 3 (clinic’s own street)	120
Adults, other wards	30
Children, Ward 3	40
Children, other wards	10
Elderly, any ward	0
Patients who speak only Kannada	4

THE TOWN AS IT REALLY IS

Group	Share of town
Adults total	55%
Children total	30%
Elderly total	15%
Ward 3 residents	18%
Kannada-only speakers	22%

Percent to write: $\text{records} \div 200 \times 100$. A share 2x or more off from reality is a red flag — mark it.

Group	% of dataset	Town share	Over / Under / Missing?
Adults, Ward 3		18%	
Children, other wards		30% total	
Elderly		15%	
Kannada-only speakers		22%	

The verdict in one sentence — who will this machine serve best, and who will it fail?

MY AUDIT HABIT — the first real place I will use percentages this month:

I will audit ... because its guest list matters to ...

WHO DECIDED? — the question behind every guest list: For the clinic dataset, one person quietly decided who got invited. Name the two people who **SHOULD** have been in that room:

WHY NOT COLLECT EVERYONE? Name one real reason a project cannot survey every person it serves — and what that forces the team to become good at:

WORD POWER • 4

Definition: sampling

Choosing part of a group to study, hoping it fairly represents the whole.

Say it three times today — it sticks!

2.2 The sampling game: small groups, big surprises

Why not just collect **everyone**? Cost, time and energy force every real project to sample — to study a part and treat it as a stand-in for the whole. The danger lives in **how** the part is chosen. Pick your sample from the people nearest to you, and the sample quietly inherits all your neighbourhood's habits. Pick it at random from the whole group, and the surprises begin to even out. You are about to feel this in one round of a game — the class meeting 8 students will **not** look like the class.

PREDICT THE GAME — BEFORE YOU PLAY IT

The first 8 hands will lean toward which kind of student? And which kind of student will the bag draw fairly? Write both predictions — we check them against the real numbers soon:

A sample I already trust (and why): _____ — one thing it still misses:

T8-05 Sampling Game

CLASS • 15 MIN

Round 1: your teacher picks the **first 8 students** who raise a hand (or the front row). Each writes one number: **minutes spent on screens yesterday**. Record them, then find the class's average for the **whole class** (your teacher writes it on the board without revealing it until Round 3).

What we measured	First-8 sample	Whole class (revealed last)
Screen minutes — average		
How far the sample average was from the truth	_____	

Round 2: now sample 8 **fairly** — assign everyone a number and draw 8 from a bag. Record the new average and the new gap:

What we measured	Fair random 8	Gap from truth
Screen minutes — average		

Which sample landed closer to the truth — and why was the hand-raising sample biased before anyone wrote a single number?

A builder collected training photos only from her own school. Which round of this game is she stuck in?

WORD POWER • 5

Definition: **representation**

How closely a dataset's people match the real people the system will serve.

Say it three times today — it sticks!

DETECTIVE DRILL · SPOT THE SKEWED SAMPLE

Each plan collects data for a school-canteen menu predictor. Mark each sampling plan F (fair) or S (skewed), and name the group it would miss:

The sampling plan	F or S?	Who is missing
Ask the 12 students in the robotics club.		
Put a slip in every lunch bag for one week; use whoever replies.		
Number all 240 students; draw 60 from a box.		
Survey only students who buy canteen food.		
Survey at 8:00 am only — bus students have not arrived yet.		

*Trickiest: the lunch-bag plan **sounds** fair — but it only hears from families that pack lunch and from students willing to answer. “Random-ish” is not random.*

2.3 Fairness Fix-It: repairing a broken guest list

Finding the skew is detective work; **fixing** it is builder work — and there is never one magic fix. Real teams choose among several repairs, each with a price. Collect **more** data from the missing groups (slow but strongest). **Re-weight** what you have (fast, but the thin groups stay thin). **Restrict** the machine to the groups it truly knows, and say so honestly on the label (safest, but less useful). Notice what is **not** on the list: pretending the machine is neutral. A machine trained on a skewed guest list is not neutral — it votes for the majority group every single time.

SAMPLE-SIZE INTUITION — BEFORE YOU CHOOSE A FIX

A re-balance can divide by ten, but it cannot invent examples. In one line: what is the DIFFERENCE between hearing a group 4 times instead of 40 times?

DETECTIVE DRILL · WHICH FIX WOULD YOU DIAL?

Each team has found its skew and must dial one fix — C (collect more), R (re-balance what exists) or L (limit the machine and label it honestly). Choose, and write the price:

The situation	C / R / L?	The price in one line
The fair-weather bus data can be collected — the monsoon is two months away.		
The exam is on Friday; there is no time to gather new examples.		
The clinic cannot reach elderly patients safely this season.		
New data exists, but only if the school pays for a second survey round.		

There is no universally right dial — a fix is right when its price is one the team can honestly pay. That is a builder's judgement, not a maths answer.

PRICE CHECK — WHAT DOES A FIX ACTUALLY COST?

The fix	Weeks of work (our estimate)	The group it helps first
Collect 40 new elderly records (home visits)		
Re-balance the existing 200 records		
Restrict to adults + print an honest label		

MY NEIGHBOURHOOD GUEST LIST — where do OUR examples come from?

If someone collected “how people cross the road” only near our school gate, which crossing habit of the wider neighbourhood would they completely miss?

T8-06 Fairness Fix-It

GROUP • 20 MIN

Return to the clinic's fever helper (T8-04): zero elderly records, 2% Kannada-only records. Your team must propose **three different fixes**, one of each kind, and be honest about the price of each:

Fix	Exactly what we would do	The price we pay (time / cost / usefulness)
1 · Collect		
2 · Re-weight or re-balance		
3 · Restrict honestly		

Which fix would you ship first, and what would you tell the ward office about what the machine still cannot do?

WORD POWER • 6

Definition: **fairness**

A system works equally well for every group of people it affects.

Say it three times today — it sticks!

■ MYTH BUSTER

MYTH “Fairness means treating everyone the same.”

TRUTH Same treatment is only fair when everyone started in the same situation. The clinic helper that scores every patient against adult, Ward-3 patterns treats everyone “the same” — and that is exactly the unfairness. Real fairness asks a harder question: **does the system work equally well for every group it touches?** Sometimes fair treatment means collecting **different** extra data for different groups — more elder-friendly examples, more Kannada-language cases — so the machine's **outcomes**, not just its rulebook, come out level. Fairness is measured in results, not in uniformity.

2.4 Consent: the door every example walks through

Behind every row of a dataset stands a person who breathed, spoke, walked or was photographed — and every honest project asks before it takes. **Consent** means permission given freely, by someone who understands what will be collected, why, and who will see it. Children and other vulnerable groups deserve extra care: often a parent or guardian must say yes alongside them. And consent has a shy cousin called **privacy care**: even data given happily must be stored safely, used only for the promised purpose, and stripped of details that could identify someone — a name, an address, a face. Builders who skip this step do not just break rules; they break the trust their whole project runs on.

CONSENT — TRUE OR FALSE?

The claim	T or F?	The one-line why
"They said yes quickly, so we do not need to explain the project."		
"Blurring faces afterwards can fix consent that was never asked."		
"Data collected for the canteen survey may be reused for attendance."		

OUR CONSENT PROMISE — PRACTICE COPY

Draft the one sentence you would say to the class BEFORE the corridor monitor collects anything. The best drafts name what, why, who sees it, and how to say no:

Self-grade your draft: names WHAT is collected ☐ · says WHY ☐ · says WHO sees it ☐ · offers a NO path ☐

WHO SEES IT? — THE FORGOTTEN CONSENT QUESTION

If the noise monitor's numbers went to...	My comfort (1–3)	What I would demand first
...only the class teacher's register		
...a wall chart everyone can read		
...the school's app for parents		

Comfort levels are data too — and they are exactly the data most projects never collect.

T8-07 Consent Check

PAIR • 10 MIN

A school AI club wants to build a “corridor noise monitor” that records sound levels to suggest quieter break times. For each data idea below, mark: G (fine to collect with plain consent), A (needs **extra** care — guardian consent, or faces/voices anonymised), or N (should not be collected at all). Then write the team’s one-sentence consent promise:

Data the club wants	G / A / N?	Why — one line
Sound level numbers (no words, no voices)		
Actual voice recordings from the corridor		
Photos of students’ faces while noisy		
Blurred silhouettes with noise numbers attached		
Names of the loudest five students, published on a wall		

Our consent promise (one sentence we would say to every student before collecting anything):

WORD POWER • 7**Definition: consent**

Permission given freely by a person who understands what will be collected and why.

Say it three times today — it sticks!

WORD POWER • 8**Definition: privacy**

Keeping people’s personal details safe, used only for the promised purpose.

Say it three times today — it sticks!

Self-check

shade one: first circle = not yet · middle = almost · last = got it

- ☐ I can audit a dataset with percentages and name the groups that are missing or overweighted ○○○
- ☐ I can explain why a hand-raised sample is biased before anyone answers a question ○○○
- ☐ I can propose three kinds of fairness fixes — and state the price of each ○○○
- ☐ I can decide what data needs plain consent, extra care, or no collection at all ○○○

■ THINK & INK

One machine or service I know was surely built on a skewed guest list. The group it was built without is ..., and one thing that changes for them is ...

PERCENT PRACTICE — FINISH THE AUDIT

The clinic audit in numbers — complete each row (records ÷ 200 × 100):

Group	Records	% of 200	Verdict
Adults, Ward 3	120		
Children, other wards	10		
Kannada-only speakers	4		

Verdict vocabulary: **overweighted** (far above the town share) · **underweighted** (far below) · **missing** (zero or nearly zero). A group cannot argue with its own percentage.

The one audit finding I will remember longest (group, number, feeling):

AT HOME Guest List Detective

HOME • 15 MIN

With a grown-up, pick one app or service your family uses (a video app, a map, a shopping site). Do not open any settings or accounts — just talk. Guess together: whose voices, faces, languages and places were probably in its training data — and who might be missing?

Our guess at the guest list (who is inside):

Who is probably missing — and one way that could go wrong for them:

One question we would ask the makers if we met them:

CASE NOTES — WHAT CHANGED IN MY THINKING?

Before this chapter I believed data was just... data — neutral stuff. Now I know every dataset has a guest list, and the detective's first question is always ...

CHAPTER 2 CLUES — POCKET THESE

- A dataset is a **guest list** — first ask who is inside, in percentages, not feelings.
- Hand-picked samples inherit the picker's neighbourhood; **random** samples land nearer the truth.
- Three honest fixes for a skewed dataset: **collect more, re-balance, or restrict and say so.**
- **Consent** is free, informed permission — and privacy means keeping the promise after it is given.

The clue I would tell my family tonight:

MY SAMPLE, AUDITED BY ME — the honest version:

My own T8-05 sample leaned toward ...

 · the fair version would need ...

Chapter 2 Checkpoint one line each — no peeking back!

ALONE • 3 MIN

1. The clinic dataset had 120 adult Ward-3 records out of 200. What percentage is that — and why is it a red flag against an 18% town share?

2. Name the three fixes for an unfair dataset, and the price of the one you would choose.

3. Why is “the machine treats everyone the same” NOT a proof of fairness?

MY CASE JOURNAL — TODAY'S SHARPEST CLUE ABOUT DATA, FAIRNESS AND CONSENT

3. How Models Learn & What No-Code Tools Do

GUIDING QUESTION

How does a machine decide — and how do we measure whether it decides well?

Now the workshop floor. You will hand-run the simplest learning machine there is — a nearest-example classifier that judges new cases by their most similar labelled examples — and then grade it with a real accuracy percentage on a hidden test set. After that you step back and survey the no-code tool market: chatbots, image classifiers and data predictors, each with inputs, outputs and hard limits. By the end you will know which tool fits which problem — and which tool marketing would love you to buy for the wrong job.

MISSIONS T8-08 – T8-10
OUTCOMES 8.L2 · 8.W1 **STRANDS**

L W

Links: Maths — distance on a grid, percentages · Science — leaf features · Social Science — honest advertising

In this chapter

3.1 The nearest-example machine	32
3.2 The honest score: accuracy on a hidden test	34
3.3 The Paper Classifier Lab: build, test, fail, improve	37

WORDS YOU'LL MEET HERE

nearest neighbour accuracy no-code tool classifier

WARM-UP · BEFORE YOU READ

ALONE · 2 MIN

Your new neighbour asks “am I like the family on the left or the right?” — how would YOU decide? Write the one feature you would look at first.

DETECTIVE'S TOOLKIT — YOU WILL NEED: pencil · ruler · calculator or rough column · 10 blank slips for test cards

3.1 The nearest-example machine

Here is the simplest honest answer to “how does a machine decide?” — **it compares**. A nearest-example classifier stores labelled examples, and when a new case arrives it finds the most similar examples it has seen and borrows their labels. That is all. No magic, no understanding — just similarity. On a grid, “similar” has a precise meaning you can measure with a ruler: the **shortest distance**. Two dots close together are two similar cases; the new point takes the label of whichever class lives nearest. You ran a cousin of this machine in Class 7; today you run the real one, with distances and a score.

RULER WARM-UP · WHICH DOT LIVES NEAREST?

A new bird sits at (5, 5). Three labelled dots live nearby: P at (4, 6), Q at (7, 5) and R at (5, 2). For each pair of dots, which is closer to the new bird — and what label would the machine vote if P and R are POND birds and Q is a FOREST bird?

Compare	Nearer to (5, 5)?	So the machine leans...
P (4, 6) vs Q (7, 5)		
Q (7, 5) vs R (5, 2)		
P (4, 6) vs R (5, 2)		

Ruler trick: count grid squares left-right and up-down, then compare. Your eyes are good at this — trust the ruler more than the vibes.

MY RULER METHOD — how I will decide “nearest” in T8-08


In one line: I will find the nearest dot by ... and if two dots are almost tied I will ...

WHICH FEATURE DECIDES? A bird sits halfway between the pond group and the forest group. Body length says FOREST; wing-beat says POND. In one line: what should the machine (and the birder) do?

T8-08 Nearest Neighbour by Hand

PAIR • 20 MIN

The plot shows a bird-watching classifier's training examples. **Teal dots** = examples labelled POND bird; **amber dots** = examples labelled FOREST bird. The x-axis is **body length (units)**, the y-axis is **wing beat rate (units)**. Three new birds arrive — the hollow numbered circles. For each, use a ruler: find the nearest teal or amber dot, and assign the label.



New bird	Nearest dot (label + its position)	My classification
1 (4.0, 5.2)		
2 (2.2, 3.0)		
3 (8.4, 4.8)		

Bird 3 is the interesting one: it sits near the middle — long body like a forest bird, low wing-beat like a pond bird. **Which feature decided its case, and what would you tell the app's users about cases like this?**

WORD POWER • 9

Definition: nearest neighbour

Classifying a new example by its most similar labelled examples.

Say it three times today — it sticks!

Definition: accuracy

The share of test examples a machine gets right, written as a percentage.

Say it three times today — it sticks!

3.2 The honest score: accuracy on a hidden test

You measured accuracy informally last year; now it gets a formula a builder can defend.

Accuracy = correct predictions ÷ total test cases × 100%. Two rules make the number honest.

First, the test set must be **hidden** during training — a machine graded on its own practice questions tells you nothing (you proved this in Class 7's Practice vs Exam). Second, one overall percentage can hide **which** cases fail — so serious builders also count failures **per group**.

Accuracy answers “how often is it right?”; it never answers “right for whom?”.

TRUE OR FALSE — ACCURACY EDITION

The claim	T or F?	The one-line why
A machine tested on its own training examples gives an honest accuracy.		
95% accuracy means the machine is right for every group it serves.		
A bigger hidden test set makes the accuracy number steadier.		
A machine can be accurate on sunny photos and useless at night.		

The two honesty rules in my own words — rule 1 (hidden test):

rule 2 (count failures per group):

PER-GROUP ACCURACY — WHAT THE ONE NUMBER HIDES

A mango scanner was tested on 100 photos: 60 sunny, 40 shaded. It scored 58 sunny hits and 8 shaded hits. Finish the arithmetic, then answer:

Group	Correct	Test cases	Group accuracy
Sunny photos	58	60	
Shaded photos	8	40	
Overall — the number the vendor will advertise	66	100	

What does the 66% hide that the vendor will never say out loud?

Which group pays the price for that advertising?

Write the honest label that belongs on the scanner's box:

THE ACCURACY CONTRACT — BUYER AND BUILDER, ONE SENTENCE EACH

The buyer will ask: “accuracy on WHICH photos, counted HOW?” — the builder will answer honestly:

T8-09 Tool Match

PAIR • 10 MIN

Match each problem to its no-code tool type — image classifier, chatbot, or data predictor — and write its likely **input**, **output**, and one hard **limit**. The first row is started for you.

The problem	Tool type	Input	Output	One hard limit
Sort mangoes as ripe or unripe at the depot	image classifier	photo of mango	ripe / unripe	needs good lighting; no reasons given
Answer parents' questions about school timings				
Predict next week's canteen rice demand				
Spot cracked solar panels from drone photos				
Guess which students may need homework help				

The trick question: a shopkeeper says “one chatbot can do all five jobs — it’s magic!” Write your two-sentence reply as the classroom’s honest tool advisor:

WORD POWER • 11

Definition: no-code tool

A ready-made AI tool you train or use without writing any program.

Say it three times today — it sticks!

DETECTIVE DRILL · WHICH TOOL — AND WHAT COULD GO WRONG?

For each promise from a (fictional) vendor, write the tool type and the sharpest thing a **Question-Habit** customer should ask before buying:

The vendor's promise	Tool type	The question I would ask
"This app looks at a leaf photo and names the disease!"		
"This chatbot answers any exam doubt, always correct!"		
"This predictor knows exactly how many idlis Sunday needs!"		

Remember the last column of the who-learns-what map: every tool's limit lives where its learning ended.

3.3 The Paper Classifier Lab: build, test, fail, improve

Time for the full experience in miniature. In one mission you will train a classifier, test it honestly, watch it fail, and fix it — the whole cycle in 25 minutes. This is the mission professional teams never skip, because **the failures are the syllabus**. Your lab: two leaf types from the school garden (or from the cards your teacher provides), a feature table, and a test set you must not peek at.

FEATURE PRE-THINK — 60 SECONDS BEFORE YOU COLLECT

The two features I plan to measure are ... and ..., because I expect them to separate the leaf types better than colour would.

LAB BENCH — WRITE YOUR ROLES BEFORE YOU TOUCH A LEAF: Collector: _____

Scribe: _____

Ruler: _____ Judge (keeps the test set honest): _____

SAFETY AT THE BENCH: collect only fallen leaves — never strip a living plant · wash hands after the garden · the Judge hides the test set BEFORE training begins, not after.

T8-10 Paper Classifier Lab

GROUP • 25 MIN

Work as a lab team of four. Roles: **Collector** (gathers leaves), **Scribe** (fills the table), **Ruler** (measures features), **Judge** (hides the test set and keeps score).

Step 1 — Collect. Gather 6 leaves of type A (e.g., neem) and 6 of type B (e.g., hibiscus). The Judge secretly keeps 2 A-leaves and 2 B-leaves aside as the **test set**. **Step 2 — Train.** For the 8 training leaves, fill the feature table and circle the feature(s) that separate the types best.

Leaf	Edge: smooth / toothed?	Length (cm)	Label A or B?
1			
2			
3			
4			
5			
6			
7			
8			

Step 3 — Write your rule. “A leaf is type B when ...”

Step 4 — Test. The Judge brings out the 4 hidden leaves. Apply your rule, then have the Judge reveal the true labels.

Hidden leaf	My rule's answer	True label	Right?
T1			
T2			
T3			
T4			

My accuracy: _____ ÷ 4 = _____ %

Step 5 — Improve. Name the one thing you would change in version 2 (more examples? a second feature? a measured threshold instead of a vague word?):

Definition: classifier

A machine that sorts examples into named groups, like ripe or unripe.

Say it three times today — it sticks!

■ MYTH BUSTER

MYTH “No-code means no thinking.”

TRUTH No-code removes the **typing**, never the **judgement**. Somebody still scopes the problem, chooses and checks the dataset, decides which examples teach the tool, picks the honest test, and reads the failures. The Plant Doctor team (Chapter 1) needed no code either — and their project still failed on data. A no-code tool in careless hands is a fast way to build the wrong thing confidently. The thinking you are practising in this book — audit, test, improve — is exactly the part no software can do for you.

Self-check

shade one: first circle = not yet · middle = almost · last = got it

- | | |
|---|-----|
| ☐ I can classify a new point by nearest labelled example — with a ruler and a reason | ○○○ |
| ☐ I can compute accuracy on a hidden test set as $\text{correct} \div \text{total} \times 100\%$ | ○○○ |
| ☐ I can match a problem to classifier, chatbot or predictor, with input, output and one limit | ○○○ |
| ☐ I can run a train → test → improve loop and name what version 2 changes | ○○○ |

■ THINK & INK

The nearest-example machine is basically “judge by closest match”. One situation where that habit would be UNFAIR to a person is ..., because ...

DETECTIVE DRILL · ACCURACY ARITHMETIC

Compute each accuracy, then answer the honesty question below:

The hidden test	Correct	Total	Accuracy %
Bird classifier, pond vs forest	17	20	
Mango ripeness scanner	94	100	
Cracked-panel spotter	9	10	

But wait — the cracked-panel spotter's 10 test cases were ALL sunny-day photos. Its real job includes night patrols. What would you tell the buyer about that 90%?

CASE NOTES — WHAT CHANGED IN MY THINKING?

Before this chapter I believed a machine decides by rules somebody typed. Now I know the nearest-example machine actually decides by ...

CHAPTER 3 CLUES — POCKET THESE

- A nearest-example classifier **compares distances** — similarity, not rules, does the deciding.
- **Accuracy** = $\text{correct} \div \text{total} \times 100\%$ — and one overall number can hide which groups fail.
- Classifier, chatbot, predictor: each has inputs, outputs and a **hard limit** where its data ran out.
- No-code removes the typing, never the thinking — audit, test and improve are still your job.

The clue I would tell my family tonight:

MY LAB REPORT CARD — THE LOOP IN ONE SMALL BOX

Version	My accuracy	The change I made, in one line
v1 (first rule)		—
v2 (after the fail)		
v2 target if we collected 8 more leaves		

Teams that write the loop down improve faster than teams that merely feel it. Three rows, three honest numbers — that is a builder's diary.

Chapter 3 Checkpoint one line each — no peeking back!

ALONE • 3 MIN

1. New point at (3.5, 6.0): nearest teal dot is at (3.0, 5.4), nearest amber dot at (6.2, 6.4). Which label — and what measurement decided it?

2. A test set has 25 cases; the machine got 21 right. Its accuracy? Show the division.

3. A vendor says his chatbot “solved” mango ripeness. What is the right tool for that job, and what can the chatbot NOT see?

MY CASE JOURNAL — TODAY'S SHARPEST CLUE ABOUT HOW MODELS LEARN

4. Ethics & Responsible AI

GUIDING QUESTION

Why does a machine that sounds so right sometimes get it so wrong — and who answers for that?

Meet the smoothest talker in the AI city: the text generator. You will play its game yourself — predicting next words from patterns — and discover why the very trick that makes it fluent also makes it **invent**. Then you hunt hallucinations in a real-looking AI paragraph, take a scenario to court as judge, lawyer and witness, and write your own AI-use pledge. This chapter has no magic answers; it has something better — a routine you can run on every confident machine you will ever meet.

MISSIONS T8-11 – T8-14
OUTCOMES 8.W2 · 8.R2 **STRANDS**

R W

Links: English — language patterns & argument · Social Science — privacy & public impact · Maths — likelihood

In this chapter

4.1 The next-word machine: how text generators actually work	43
4.2 Hallucination Hunt: verify before you trust	44
4.3 The Responsible-Use Court: decisions with consequences	46
4.4 My AI-Use Pledge	47

WORDS YOU'LL MEET HERE

text generator

hallucination

misinformation

responsible use

WARM-UP · BEFORE YOU READ

ALONE · 2 MIN

Finish this sentence the way a friend probably would: “The early bird catches the ...”. Now finish: “The early crow catches the ...”. Which one were you surer about — and where did that knowledge come from?

DETECTIVE'S TOOLKIT — YOU WILL NEED: pencil · the printed paragraph in T8-12 · sticky notes for court roles · an honest heart

4.1 The next-word machine: how text generators actually work

Forget magic. A text generator does one thing, enormously well: it reads the words so far and predicts a **likely next word** — based on patterns in a mountain of human writing. It is autocomplete that grew up. The catch is built into the trick itself: the machine knows which words **usually follow**, not which words are **true**. “The early bird catches the worm” flows because millions of people wrote it; a machine completing “The early crow catches the...” will pick whatever sounds plausible. Usually it sounds right. Sometimes it is completely invented — with perfect confidence either way, because confidence and truth live in different rooms.

T8-11 Next-Word Game

CLASS • 15 MIN

You are the machine. Your teacher reads each sentence aloud up to the blank; your class votes on the most likely next word. Then your teacher reveals the counts from a 500-page book of local writing (invented for this game). Fill the table as you play:

Sentence up to the blank	Class's first guess	Runner-up guess	Revealed: which word won, and by how much
“The monsoon arrived early, bringing heavy ...”			
“She opened the lunch box and found her favourite ...”			
“The cricket match was delayed because of ...”			
“The old elephant walked slowly toward the ...”			

Now the deep question: when the machine picked a word nobody in the room expected, was it **wrong**? What was it actually optimising for?

A classmate asks: “if it’s just picking likely words, why do its essays sound so intelligent?” Answer in one line:

Definition: **text generator**

A machine that writes by predicting the most likely next words.

Say it three times today — it sticks!

DETECTIVE'S NOTE

Here is why the machine **sounds** smart: likely words, arranged by likely patterns, produce paragraphs shaped exactly like the ones people write. It is a brilliant imitation of the **shape** of knowledge. But shape is not truth. A hallucinated biography of a fake scientist has perfect grammar, a plausible structure and a confident tone — all borrowed from real writing about real scientists. That is why your eyes and your verification routine are still the most important parts of the system.

4.2 Hallucination Hunt: verify before you trust

When a text generator states something false as if it were fact, builders call it a **hallucination** — and it is not a bug that will be patched away next month; it is the direct result of predicting likely words instead of checking facts. So the responsible user carries a routine, not a hope. The Class 6 habit — **check before believing** — now grows up into three moves: **trace the claim** (is there a real source?), **test the number** (do two sources agree?), and **smell the story** (is it suspiciously exactly what I wanted to hear?).

SAME SHAPE, DIFFERENT TRUTH — A PRE-HUNT DRILL

Both sentences below are grammatical, confident and plausible. Only one survives a fact-check. Mark which — and name the exact words that made the other one suspicious:

The sentence	V (verifiable) or S (suspect)?	The words that gave it away
"The Nagpur district library opens at 8 am on weekdays."		
"Most people agree the town's founder was a poet-king in 1238."		

Watch the give-away words: **most people agree**, **everyone knows**, **it is said that** — confidence borrowed from a crowd nobody can name.

Below is a paragraph a chatbot produced when a student asked about their town. Hunt it: mark each claim V (verifiable and checks out), F (false — contradicted by the fact box), or U (unverifiable — no way to check from here). Then write the chatbot a revision order.

WHAT THE CHATBOT WROTE

“Sunapur was founded in 1847 by the poet R. Krishnamurthy, who also invented the sunflower oil press used across the district today. The town’s famous three-day Mango Festival, started in 1952, attracts over 5 lakh visitors every year. The Government Higher Primary School, established 1961, was the first school in the state to teach computer science. The town bus route 14 runs every ten minutes from the railway station to the fort.”

Claim	V / F / U
Founded in 1847 by R. Krishnamurthy	
He invented the sunflower oil press	
Mango Festival began 1952, 5 lakh visitors	
School was FIRST in the state to teach computing	
Bus route 14 every ten minutes to the fort	

THE FACT BOX (from the town library)

· Sunapur was founded in **1884**; its founder was a merchant, not a poet. · The oil press story is **unknown to any local record**. · The Mango Festival began in **1972**; visitor numbers are **not published anywhere**. · The school dates from 1961 — true — but the first state school to teach computing was in another district. · Route 14 runs **hourly**, and the town has no fort.

Notice: every invented claim was **shaped** like a true one — a date, a founder, a festival. The machine matched the pattern, not the facts.

My revision order to the chatbot (which claims to fix, and what evidence I demand):

WORD POWER • 14**Definition: hallucination**

When a text generator states something false as if it were fact.

Say it three times today — it sticks!

Definition: misinformation

False content that spreads — by mistake or on purpose.

Say it three times today — it sticks!

DETECTIVE DRILL • THE THREE-MOVE ROUTINE

Name the move: T (trace the claim), N (test the number) or S (smell the story):

What the responsible user does...	Which move?
Searches for a second, independent source that says the same thing.	
Asks: “who gains if I believe this instantly?”	
Opens the original article the chatbot claims to summarise.	
Checks whether “92% of doctors agree” names WHO counted and HOW many.	
Feels a thrill of confirmation — and slows down on purpose.	

■ MYTH BUSTER

MYTH “If it is trained, it is true.”

TRUTH Trained means the machine found patterns in its examples — nothing more. The Plant Doctor was trained and still failed river-farm tomatoes; text generators are trained on human writing and still invent facts, because their job was never truth-finding, it was pattern-continuing. Training makes a machine **consistent** with its examples. Only **testing against the world** — hidden test sets for classifiers, source-checking for text — makes it trustworthy. Whenever you hear “our AI is trained on millions of examples”, respond like a builder: **trained, yes — but tested how, and on whom?**

4.3 The Responsible-Use Court: decisions with consequences

Technology choices are people choices. A school, a company or a government decides **what** a machine may do, **to whom**, and **who answers when it fails**. These decisions deserve the same seriousness as any other — evidence, arguments, and a verdict someone signs. So: court is now in session.

The case: Sunapur school installed a “focus camera” in Class 8-B that watches faces and flags students who “look distracted” to the teacher every 10 minutes. The vendor says it raises results. Two students say it feels like being suspected all day. One parent says her son’s medication makes him look drowsy — the machine flags him constantly.

Assign roles: one **Judge**, two **Defence** lawyers (for the camera), two **Prosecution** lawyers (against), two **Witnesses** (the parent; a Class 8-B student), and the rest are the **Jury**. Each side gets 3 minutes; witnesses 1 minute each. Then everyone fills the verdict sheet:

Court question	The defence's best point	The prosecution's best point
Does the machine measure what “focus” really is?		
Who is harmed if it is wrong — and how?		
Was consent taken from the people watched?		
Is there a less invasive way to the same goal?		

THE JURY’S VERDICT (pick one and give the reason): keep the camera · keep with changes (name them) · remove it.

WORD POWER • 16

Definition: responsible use

Choosing and using AI tools so people are helped, respected and safe.

Say it three times today — it sticks!

4.4 My AI-Use Pledge

Court cases end with judgements; builder stories end with promises. A pledge is not decoration — it is the checklist you write **before** the difficult moment arrives, when a chatbot offers to write your essay or a group-chat spreads an exciting rumination. Copy your class’s best ideas into the box, sign it, and keep the page where you will meet it again at the end of the book.

Complete each clause honestly — these are YOUR rules, not the teacher's:

I, the undersigned investigator, pledge that...

1. Before I believe or forward any AI-written or AI-edited content, I will

2. When a text generator helps me with schoolwork, I will

3. When I am tempted to feed personal information — mine or anyone's — into an AI tool, I will

4. When an AI tool makes a decision about a person, I will ask the two court questions:

5. If a machine's output hurts someone, I will hold responsible the ... (circle one: machine / its user / its maker / whoever had control) — **because**

Signature: _____ Date: _____

Self-check

shade one: first circle = not yet • middle = almost • last = got it

- | | |
|---|-------|
| ☐ I can explain in two sentences how a text generator writes — and why it can invent | ○ ○ ○ |
| ☐ I can run the three-move routine — trace, test, smell — on a suspicious claim | ○ ○ ○ |
| ☐ I can argue both sides of a responsible-use case and sign a reasoned verdict | ○ ○ ○ |
| ☐ I can name what my own pledge demands of me before the difficult moment arrives | ○ ○ ○ |

■ THINK & INK

A chatbot once confidently told me something that turned out to be wrong (or would have, if I had believed it). Now I understand why that happened: ...

HOME LINK take this page home — no screen needed

AT HOME Family Fact Patrol

HOME • 15 MIN

With a grown-up, find one forwarded message or AI-looking summary in the family phone (do not open links). Run the routine together: **trace** — where did it really come from? **test** — do two trusted sources agree? **smell** — does it push an emotion before evidence?

The claim we examined:

What the three moves found:

What we will do differently before forwarding next time:

CASE NOTES — WHAT CHANGED IN MY THINKING?

Before this chapter I believed a confident, fluent answer was usually a correct answer. Now I know fluency and truth are built differently, because ...

CHAPTER 4 CLUES — POCKET THESE

- A text generator **predicts likely words** — likely is not the same as true.
- **Hallucinations** are built from the same fluency as facts — verify with trace, test, smell.
- Responsible use is decided by **people** — in courtrooms, staffrooms and pledges, with reasons.
- Responsibility follows **control**: machines control nothing, so people answer — makers, users, deployers.

The clue I would tell my family tonight:

MY COURT RECORD — THE VERDICT I SIGNED

Case: the focus camera. I argued for the side of ... · my verdict: ...
one argument that changed my mind mid-debate:

Chapter 4 Checkpoint

one line each — no peeking back!

ALONE · 3 MIN

1. Why does a text generator hallucinate — what exactly is it doing instead of checking facts?

2. A forwarded message says “92% of parents demanded a focus camera”. Which of the three moves do you run first, and what do you look for?

3. The court’s hardest question: the camera’s **maker**, the school that **installed** it, and the teacher who **uses** it — who answers when it wrongly flags a student? Why?

MY CASE JOURNAL — TODAY’S SHARPEST CLUE ABOUT ETHICS AND RESPONSIBILITY

5. Capstone: The AI Project Proposal

GUIDING QUESTION

Can you plan an AI that people can trust — and defend the plan to a live audience?

The final case brings every clue together. Your team will write a complete AI project proposal on paper — sharp problem, audited data plan, chosen tool, honest test, fairness check and known limits — exactly the document a professional team brings to its first review. Then your classmates will pressure-test it in a peer review gallery, and you will sharpen your budget choices with a real optimisation problem. This is not an exam; it is the moment the detective becomes the builder.

MISSIONS **T8-15 – T8-17**

OUTCOMES **8.U1 • 8.L1 • 8.L2 • 8.D1 • 8.D2 • 8.W1 • 8.W2 • 8.R1 • 8.R2 • 8.T1**

STRANDS

U D L W R

Links: all five strands • English — persuasive writing • Maths — budgeting & optimisation

In this chapter

5.1 The proposal: your plan on paper	52
5.2 The Peer Review Gallery	54
5.3 Optimise It: the best you can afford	56

WORDS YOU'LL MEET HERE

peer review

trade-off

optimise

stakeholder

WARM-UP • BEFORE YOU READ

ALONE • 2 MIN

You have one page and one minute to convince your class to build your AI. What is the FIRST thing you would say — the problem, the data, or the dream? Write it, then check if your answer survives Task T8-15.

DETECTIVE'S TOOLKIT — YOU WILL NEED: your whole notebook • pencil AND eraser • the fairness checklist from Chapter 2 • courage to share

5.1 The proposal: your plan on paper

Every serious AI project begins as a document nobody coded: a proposal. Its job is to answer six questions **before** anyone builds — What problem, exactly? What data, from whom, with what consent? Which tool, and why that one? How will we test it honestly? Who could it hurt, and what did we change for them? What can it NOT do? A proposal that answers all six can still fail — but a project without one fails **surprised**, which is worse. Your proposal lives in Tasks T8-15 over the next pages. Choose a real, local, small problem: your classroom, your street, your school's morning chaos. Small and real beats grand and imaginary every time.

T8-15 • Part 1 AI Project Proposal — the plan

GROUP • 3 periods MIN

Part 1 of 2. Work as a project team of three or four. Choose your problem, then scope it with the discipline of Chapter 1.

PROPOSAL • SECTION A — THE SHARP PROBLEM

Our problem statement (passes the 4W test):

The stakeholders — everyone this machine would touch, helped or harmed (name at least three groups):

Why existing solutions are not enough:

SECTION B — THE TOOL

Which no-code tool type fits — classifier, chatbot or predictor — and why that one:

Its input, its output, and one hard limit we accept:

T8-15 • Part 2 AI Project Proposal — data, test & fairness

GROUP • 3 periods MIN

Part 2 of 2. This is where the Plant Doctor failed. Do not skip a box.

SECTION C — THE DATA PLAN

What we collect	From whom (all the groups!)	How we get consent	How we protect privacy

Our guest-list check (from Chapter 2): which groups must be INSIDE the data for this to be fair — and how many from each?

SECTION D — THE HONEST TEST

Our hidden test set: how many cases, and why that number:

The accuracy we would accept before shipping (and the accuracy that would send us back to the drawing board):

SECTION E — FAIRNESS & LIMITS

The group most likely to be failed by version 1 (be specific):

What we changed in the plan because of that group:

What our machine will NEVER do (write it down before someone assumes otherwise):

5.2 The Peer Review Gallery

Real proposals are never shipped on the builder's word alone — they face review. In professional teams, reviewers hunt for the **weakest box**, not the prettiest sentence, and the proposal gets better because of it. Today your classmates are your reviewers: every plan is pinned up (or read aloud), every team walks the gallery with a clipboard, and every review ends with one strength, one gap, and one question the builders cannot dodge.

GALLERY MARSHAL — HOUSE RULES: review the PLAN, never the people · one strength before any gap · every gap must come with the question that would close it · builders may only reply with evidence, not feelings.

THE WEAKEST BOX IN AI HISTORY — a 30-second recall: The Plant Doctor proposal would have scored well on THREE of the six boxes. Which box would you have graded Beginning — and what ONE line would have saved it?

MY GALLERY CLIPBOARD — plan the walk before you walk it

The proposals I will review (numbers): _____ The question I am proudest to ask:

Reviewer's oath: "I will find the weakest box, say what is strong first, and leave every plan better than I found it."

Walk the gallery. For each proposal except your own, fill one review card. Be honest and kind — you are reviewing the **plan**, never the people:

REVIEW CARD — FOR PROPOSAL N° _____

The strength I would praise first:

The weakest box right now (problem / data / test / fairness):

The one question the builders cannot dodge:

Rubric shade for this proposal — Beginning / Growing / Secure / Shining (shade one)

○ ○ ○ ○

Now the reverse mirror: the sharpest question a reviewer asked about YOUR plan, and what version 2 will do about it:

CARD 2 OF 3 — THE ONE-MINUTE VERSION (use for a third proposal)

Proposal N° _____ · weakest box: _____ · rubric shade: ○ ○ ○ ○

My closing question for the builders:

THE FOUR LEVELS — HOW THIS BOOK GRADES REASONING

Level	What it looks like in a proposal
Beginning	The problem is a slogan; the data is “we will find some”; fairness is one word.
Growing	Real problem and tool chosen, but the test is vague or the fairness check is a promise, not a plan.
Secure	All six boxes are answered with evidence — and the limits are written before anyone asked.
Shining	Secure, PLUS the team found a trade-off nobody noticed and made a defensible choice anyway.

Notice: a flawed answer with honest reasoning can outscore a polished answer with no reasons. This rubric rewards exactly what builders get paid for.

5.3 Optimise It: the best you can afford

One last thinking-tool before graduation. Real builders never have enough budget for everything — they **optimise**: they break the choice into parts (cost, benefit, fairness) and find the best combination under the limit. This is computational thinking in its purest form: no AI at all, just you, a table and a constraint.

THE OPTIMISER'S HABIT

Break the choice into parts (this item costs this, buys that), compute honestly, then defend what you leave out. The defence matters as much as the choice — a budget nobody can argue with is a budget nobody checked.

BUDGET ARITHMETIC WARM-UP — BEFORE YOU SHOP

Add up these three starter combinations. Which ones fit inside ₹500 — and which item appears in BOTH affordable pairs?

Combination	Total cost	Fits ₹500?
A + B (trip + test cards)		
C + D (two features + lighting)		
B + C + D (tests + features + lighting)		

T8-17 Optimise It

ALONE • 10 MIN

Your club has ₹500 to improve the leaf classifier before its second test. Choose the combination that gets the most trustworthy machine. You may buy each item at most once. Warm-up totals: $A+B = ₹300$ · $C+D = ₹400$ · $B+C+D = ₹500$ — all three fit, but none fixes everything. Now choose:

Upgrade item	Cost	What it buys you
A · Second field trip for 20 more leaves, half from the river farm	₹200	fixes the missing-group problem in training
B · Print 10 extra hidden test cards	₹100	bigger test set — a steadier accuracy number
C · Measure both leaf features (edge AND length) for all cards	₹150	two features instead of one — fewer mix-ups
D · Photograph everything in shade AND sun	₹250	kills the lighting weakness the Plant Doctor had

My combination (letters + total cost): _____ The accuracy problems this combination fixes:

The weakness I knowingly leave unfixed — and my defence for leaving it:

Could any combination fix everything at once? What does that tell you about real projects?

■ MYTH BUSTER

MYTH “The model knows why.”

TRUTH Ask the leaf classifier **why** it called a leaf type B, and it has no answer — only distances. Ask the text generator why it wrote a sentence, and it will **invent** a plausible reason, the same way it writes everything else. The “why” lives in three places, and none of them is the model: in the **data** (what patterns existed), in the **design choices** (which features, which examples), and in the **people** (who decided what it was for). That is why every chapter of this book made you write your reasons beside your answers — the reasoning was never the machine’s job. It was, and is, yours.

Self-check

shade one: first circle = not yet · middle = almost · last = got it

- | | |
|--|-----|
| ☐ I can write a proposal that answers all six builder questions with evidence | ○○○ |
| ☐ I can review a classmate’s plan with one strength, one gap and one unavoidable question | ○○○ |
| ☐ I can grade a proposal — mine or a classmate’s — with the four levels and say why | ○○○ |
| ☐ I can choose the best combination under a budget and defend the trade-off I accept | ○○○ |

■ THINK & INK

After this year, my honest definition of “trustworthy AI” is ... — and one thing I now check that I never checked before is ...

CASE NOTES — WHAT CHANGED IN MY THINKING?

Before this year I believed AI was mainly about clever machines. After building (on paper) my own, I believe it is mainly about ...

CHAPTER 5 CLUES — POCKET THESE

- A proposal answers six questions **before** building: problem, data, tool, test, fairness, limits.
- **Peer review** hunts the weakest box — the plan that survives honest reviewers earns its trust.
- **Optimising** means choosing the best combination under limits — and defending what you leave out.
- The model never knows why — the reasons live in the data, the design and the people.

The clue I would tell my family tonight:

Chapter 5 Checkpoint one line each — no peeking back!

ALONE • 3 MIN

1. Name the six sections of your proposal, in order — no peeking back.

2. A reviewer shades “Growing” on a proposal whose boxes are all filled. What is most likely missing?

3. In T8-17, why is “fix everything” impossible — and what does a real team do about that?

MY CASE JOURNAL — THE LAST CLUE OF CASE 3, FOR THE ROAD TO CLASS 9

My AI words

The sixteen words on this case — each one met through an experience, never before.

1 problem statement

One precise sentence saying exactly what problem you solve, for whom, where.
USE IT IN MY OWN SENTENCE: _____

2 AI project cycle

The loop: scope the problem, collect data, build and test, reflect and improve.
USE IT IN MY OWN SENTENCE: _____

3 dataset

The collection of examples a machine learns from, or is tested on.
USE IT IN MY OWN SENTENCE: _____

4 sampling

Choosing part of a group to study, hoping it fairly represents the whole.
USE IT IN MY OWN SENTENCE: _____

5 representation

How closely a dataset's people match the real people the system will serve.
USE IT IN MY OWN SENTENCE: _____

6 fairness

A system works equally well for every group of people it affects.
USE IT IN MY OWN SENTENCE: _____

7 consent

Permission given freely by a person who understands what will be collected and why.
USE IT IN MY OWN SENTENCE: _____

8 privacy

Keeping people's personal details safe, used only for the promised purpose.
USE IT IN MY OWN SENTENCE: _____

9 nearest neighbour

Classifying a new example by its most similar labelled examples.
USE IT IN MY OWN SENTENCE: _____

10 accuracy

The share of test examples a machine gets right, written as a percentage.
USE IT IN MY OWN SENTENCE: _____

11 no-code tool

A ready-made AI tool you train or use without writing any program.
USE IT IN MY OWN SENTENCE: _____

12 classifier

A machine that sorts examples into named groups, like ripe or unripe.
USE IT IN MY OWN SENTENCE: _____

13 text generator

A machine that writes by predicting the most likely next words.
USE IT IN MY OWN SENTENCE: _____

14 hallucination

When a text generator states something false as if it were fact.
USE IT IN MY OWN SENTENCE: _____

15 misinformation

False content that spreads — by mistake or on purpose.
USE IT IN MY OWN SENTENCE: _____

16 responsible use

Choosing and using AI tools so people are helped, respected and safe.
USE IT IN MY OWN SENTENCE: _____

BUILDER'S OATH

Machines **compare, predict, continue** — people **scope the problem, audit the guest list, run the honest test and sign the verdict**. When a machine speaks, I will ask my five questions: **How do I know? What is missing? Who made this? Who is missing? What is the trick here?**

ANSWER HINTS — FOR CHECKING, NEVER FOR COPYING

T8-04: Ward-3 adults are $120 \div 200 = 60\%$ of the dataset against an 18% town share — overweighted more than three times; elderly are **missing** (0%); Kannada-only speakers are 2% against 22% — nearly missing. · **T8-05:** the hand-raised sample is biased before any number is written — quick hands, front rows and confident friends are not a random draw. · **Name the stage drill:** P, D, M, T, I, P. · **Sharpen or spoil:** S, V, S, V — the vague ones are missing all four Ws at once. · **Skewed sample drill:** S, S, F, S, S — the lunch-bag plan only hears from packers and repliers.

MORE HINTS — FOR THE STICKY MISSIONS

T8-08: Bird 1 and Bird 2 both fall among the teal dots — POND. Bird 3's nearest dot is the amber one at about (8.5, 6.2) — FOREST — but write down which **feature** actually did the deciding, because the vibes disagreed. · **T8-09:** timings → chatbot; rice demand and homework-help flags → predictor; cracked panels → image classifier. · **Three-move drill:** N, S, T, N, S. · **Accuracy drill:** 85%, 94%, 90% — then ask what the sunny-day-only test set hides. · **T8-17:** every four-item combination costs ₹700, so **no** plan fixes everything — that is the lesson. Strong plans include A (missing groups) and then defend what they left out.

MY HINT LOG — WHICH MISSIONS NEEDED A PEEK?

Hints are training data for you. Tick the mission code each time a hint unstuck you — a mission that needed three hints deserves one more retry next week, not three more hints:

T8-__ needed a hint

T8-__ needed a hint

T8-__ needed a hint

T8-__ needed a hint

A POSTCARD TO THE CLASS 9 INVESTIGATOR — WHO WILL BE YOU

The next case asks the question every builder must face sooner or later: “**Why does AI work — and when does it fail?**” Class 9 goes under the hood: the maths of averages, probability and best-fit lines that make learning possible, and the graph tricks that make it look easier than it is. Leave three clues for your future self — the trade-off that surprised you most, the plan of yours that failed hardest, and one question you dare Class 9 to answer:

My progress

Shade honestly — investigators never fake evidence.

MY OUTCOME TRACKER — TICK WHEN YOU TRULY CAN DO IT

- | | |
|---|--|
| ☐ 8.U1 — define AI and machine learning in my own words; place chatbots, classifiers and recommenders on the who-learns-what map | ☐ 8.W1 — match no-code tool types to problems, stating inputs, outputs and limits |
| ☐ 8.L1 — walk a local problem through the full AI project cycle, on paper | ☐ 8.W2 — explain why a text generator can sound right yet be wrong — and verify its claims |
| ☐ 8.L2 — hand-run a nearest-example classifier, compute its accuracy, and suggest an improvement | ☐ 8.R1 — apply a fairness checklist and propose fixes for an unfair dataset |
| ☐ 8.D1 — audit a small dataset for missing groups and compute percentages per group | ☐ 8.R2 — make and defend a responsible-use decision about privacy, misinformation or impact |
| ☐ 8.D2 — explain why collecting data needs consent and privacy care | ☐ 8.T1 — break a multi-variable problem into parts and find the best choice under limits |

MY CASE FILES — SHADE A STAR PER CHAPTER FINISHED

1 Project Cycle · 2 Data & Fairness · 3 How Models Learn · 4 Responsible AI · 5 Proposal



PRATIMAI · AI HANDOUTS · CLASS 8

My Builder's Certificate

This certifies that project lead

INVESTIGATOR NAME

has completed the case file **“Build the AI Cycle”** — Level 3 · BUILD THE CYCLE — planned and peer-reviewed a full AI project proposal, and promised to keep asking: **How do I know? · What is missing? · Who made this? · Who is missing? · What is the trick here?**

DATE

INVESTIGATOR'S SIGNATURE

GUIDE'S SIGNATURE

My series so far — three cases, one investigator

CLASS 6 • NOTICE

I learned to **spot** learning machines and to ask: how do I know?

CLASS 7 • SORT & PREDICT

I learned to **train and test** machines — and how examples fool them.

CLASS 8 • BUILD THE CYCLE

I learned to **plan, audit and defend** an AI before it is built.

The habit I am proudest of now:

Where I used my five questions — the tally of a finished case

My question	Where I asked it for real (at home, in class, online)
How do I know?	
What is missing?	
Who made this?	
Who is missing?	
What is the trick here?	

Five rows is the whole method. If a row is still empty, the next case (Class 9) is where you will catch up.

THIS CASE FILE BELONGS TO THE EVIDENCE

Investigator: _____ Class & division: _____ School year: _____

Hand this book to next year's Class 8 detectives with your tallies inside — the best hand-me-down is an honest one.

■ THE DOOR TO CLASS 9

Next year the case goes **under the hood**: the averages, probability and best-fit lines that make machine learning possible — and the graph tricks that make it look easier than it is. Keep this book; its checkpoints are your revision notes. And whatever machine you meet next, keep asking your five questions:

How do I know? • What is missing? • Who made this? • Who is missing? • What is the trick here?

Ready for the next case? Class 9 asks: “**Why does AI work — and when does it fail?**” — bring your blueprint.