PRATIMA AI · CURRICULUM DESIGN

FOR SCHOOLS · CLASSES 6 TO 10

The AI Curriculum

A five-year journey — from noticing smart machines to designing them responsibly

CLASS 6	CLASS 7	CLASS 8	CLASS 9	CLASS 10
NOTICE	**PREDICT**	**BUILD**	**REASON**	**DESIGN**
the Detective	the Investigator	the Engineer	the Analyst	the Designer

PRATIMA AI Curriculum Team

96 paper missions · 25 chapters · 5 student roles · no screens needed

2026 EDITION · PREPARED FOR SCHOOL LEADERSHIP AND ACADEMIC COORDINATORS

Preface

A note from the PRATIMA AI Curriculum Team on how this five-year course is designed — and why.

WHY WE BUILT THIS CURRICULUM

Every child already lives with artificial intelligence. It picks the next video, reroutes the school bus around traffic, and answers the question a cousin whispers to a phone. What most children never get is the chance to look underneath — to see that a "smart" machine is following steps or learning from examples, that its knowledge comes from data, and that it can be wrong in ways adults miss. We built this curriculum to give every child that chance, from Class 6 to Class 10. The design follows what we know about how children learn: concrete before abstract, doing before defining, talking before testing. So nothing here starts with a definition; everything starts with something a child can hold, sort, measure or argue about.

WHAT SCHOOLS ARE USUALLY OFFERED

Schools looking for AI education usually meet one of three things. First, a textbook of definitions to memorise — accurate, perhaps, but inert for a twelve-year-old. Second, a coding course that needs a laboratory, licensed software and screens many classrooms cannot spare, which quietly divides children into those who have and those who do not. Third, material adapted from adult corporate training and dressed down with cartoons. All three share one flaw: they teach **about** machines instead of teaching children to **think** — carefully, honestly and with their own hands — about what those machines really do.

HOW WE ARE DIFFERENT

Our answer is a five-year journey built on five commitments. **Paper first:** all 96 missions run on pencil, paper and conversation — no laboratory, no screens, no child left behind by cost. **One role per year:** the Detective who notices, the Investigator who operates, the Engineer who plans the cycle, the Analyst who reasons with the mathematics, and the Designer who evaluates and designs responsibly. **Human versus machine, side by side:** students write a paragraph and compare it with one a machine wrote, judge contests between people and programs, and discover where human judgment genuinely beats pattern-matching — and where it does not. **Ethics woven through:** fairness, honesty, privacy and accountability appear in every year, because children meet AI's failures as often as its magic. **Rooted in this country:** kirana stores, monsoon data and neighbourhood surveys — the examples are the students' own world.

OUR ACTIVITIES AND TASKS

Every chapter runs on missions — structured written activities a pair or team completes in class: running a human robot, auditing an unfair dataset, holding a courtroom trial for an AI mistake, dismantling a "99% accurate" claim. Each year closes with a capstone: a neighbourhood investigation, a class-built prediction machine, a defended project proposal, an SDG-linked brief, a responsible-design pitch to a board of judges. Assessment is formative throughout — feedback in the margins, a badge earned each year, a portfolio by Class 10 — because our goal is confidence, not examination fear.

WHERE OUR STUDENTS END UP

A student who walks the full journey can do something rare: notice the systems around them, operate their core ideas with their own hands, plan an AI project the way professionals do, reason with the mathematics underneath, and finally judge and design such systems responsibly — defending every choice with evidence. They will not all become AI engineers, and they should not need to. But every one of them will study, work and vote in a world full of these machines. We want them meeting that world awake: curious, careful and unafraid to ask how it really works. That is the end goal — and every mission in the five books was built to serve it.

THE AI CURRICULUM · CLASS 6–10 · FIVE BOOKS, ONE JOURNEY — *The PRATIMA AI Curriculum Team*

CLASS 6

LEVEL 1 — NOTICE

AI Around Me

20 PAPER MISSIONS · 5 CHAPTERS

My detective notebook for spotting smart machines

YEAR 1 OF 5

Class 6 opens the five-year journey by training attention — the quiet skill underneath all AI literacy. Students notice that "smart" machines are everywhere, and learn the detective's first secret: every clever machine is doing one of two things. It is either following exact steps somebody wrote, or it is learning patterns from examples. Twenty paper missions later, they have trained machines themselves, fed learners good and bad data, chased patterns through games, guarded their privacy online, and closed the year with a real investigation of the AI in their own neighbourhood. Nothing needs a screen; everything needs curiosity. By the end of the year, a child no longer says "the computer is so smart" — the child asks: what did it learn from, and who showed it that?

5	20	14	0
CHAPTERS	PAPER MISSIONS	WORDS TO KEEP	SCREENS NEEDED

CH 1 Machines That Seem Smart

T6-01 – T6-05

When is a machine actually learning — and when is it only following steps?

OUTCOMES

- Separate rule-followers from learners through the Human Robot game and flashcard teaching
- Judge a fair humans-versus-machines contest and defend the verdict
- Use the book's three detective questions on machines they meet every day

CORE IDEAS

algorithm · automation · artificial intelligence · training examples · predict

Core understanding — Every "smart" machine is either following exact written steps or learning patterns from examples — and telling the two apart is the first skill of AI literacy.

SUBJECTS

Maths — patterns · English — exact instructions

CH 2 Data: The Food of AI

T6-06 – T6-09

What do machines learn from — and what happens when the food is bad?

OUTCOMES

- Hunt the four kinds of data hiding around the classroom
- Tidy observations into tables, and turn tables into bar charts
- Feed a learner bad examples and watch "garbage in, garbage out" happen

CORE IDEAS

data · pixel · tally · garbage in, garbage out · bar chart

Core understanding — Machines learn from data the way children learn from teachers — poor data teaches poor lessons, and a learner is only as good as its food.

SUBJECTS

Maths — tables & bar graphs · Science — observation

SIGNATURE MISSION T6-01 · Human Robot

One student becomes the Robot, the other the Engineer. The Robot must follow the written steps **literally** — and the funniest failure lands the lesson no lecture could: a machine never guesses what you **meant**, only what you **said**.

CH 3 Patterns & Decisions

T6-10 – T6-13

How does a machine decide what to do next?

OUTCOMES

- Hunt number and letter patterns, then play Guess My Rule with a hidden rule-keeper
- Trace an if-then decision machine step by step
- Challenge answers whose confidence is higher than their accuracy

CORE IDEAS

pattern · prediction · confidence · if-then · question tree

Core understanding — A prediction is a pattern turned into a decision — and every prediction carries a confidence that must be questioned, not trusted.

SUBJECTS

Maths — number patterns · English — clear questions

CH 4 Be a Smart Digital Citizen

T6-14 – T6-17

How do I enjoy the online world — and stay safe, kind and honest?

OUTCOMES

- Sort personal information with the privacy traffic light
- Cook unguessable passphrases and run the three-check real-or-fake detector
- Follow one drawing's digital footprint across the internet

CORE IDEAS

privacy · personal information · passphrase · verify · digital footprint

Core understanding — The online world runs on the same honesty and care as the offline one — what you share, whom you trust and what you forward are your choices.

SUBJECTS

Social Science — community life · English — checking facts

CH 5 My Neighbourhood AI Investigation CAPSTONE

T6-18 – T6-20

What can I discover about AI near me — using my own data?

OUTCOMES

- Plan and run a real survey of five grown-ups about the machines around them
- Organise and chart the collected data, then design a helper machine of their own
- Close the year's big case by answering page one's question — with evidence

CORE IDEAS

survey · tally · bar chart · limit · evidence

Core understanding — Evidence beats impression: a careful survey of their own neighbourhood shows children that AI knowledge begins with their own questions.

SUBJECTS

Social Science — community survey · Maths — data & charts

BY THE END OF THIS YEAR

- Tell rule-followers from learners, and explain what data taught a machine
- Read and draw charts, question a confident claim, and guard personal information
- Run a small survey and defend a conclusion with evidence

1 BADGE EARNED · LEVEL 1 — NOTICE COMPLETE

CROSS-CURRICULAR FOOTPRINT

How many chapters reach each subject



CLASS 7

LEVEL 2 — SORT & PREDICT

How Machines Learn

17 PAPER MISSIONS · 5 CHAPTERS

My case file for training machines — and catching their mistakes

YEAR 2 OF 5

In Class 7 the investigator takes over the machine. Last year students spotted learning machines; this year they operate one. They run all three learning jobs with their own hands — classification, regression and clustering — and sit the same hidden exam a machine sits, so that "practice versus exam" stops being a metaphor. They meet the machines that see, read and recommend, tour five sectors where AI already works across India, and then study how these systems fail: biased data, faked content, and the question of who answers when something goes wrong. The capstone turns the class itself into a prediction machine — built from a real survey, checked for bias, and graded against another section's test set.

5	17	14	0
CHAPTERS	PAPER MISSIONS	WORDS TO KEEP	SCREENS NEEDED

CH 1 The Learning Machine

T7-01 – T7-05

How can a machine learn without being told the rules?

OUTCOMES

- Run classification, regression and clustering by hand with real objects and numbers
- Sit the same hidden exam a machine sits — training set versus test set
- Explain in one sentence why a machine needs examples instead of rules

CORE IDEAS

classification · regression · clustering · training set · test set

Core understanding — A machine can be taught by examples instead of rules — but it must then be examined on examples it has never seen.

SUBJECTS

Maths — graphs & averages · English — precise definitions

CH 2 Machines That See, Read and Recommend

T7-06 – T7-09

What does a machine actually see when it looks at a picture, a message or you?

OUTCOMES

- Be the vision machine: turn pixels into features and predict what they show
- Be the reading machine: count words and frequencies — and meet the idiom that defeats it
- Be the recommender: find people whose pattern of likes matches yours

CORE IDEAS

feature · frequency · recommendation · context · idiom

Core understanding — A machine sees numbers, not meaning; counting is powerful, but understanding context is still a human skill.

SUBJECTS

Maths — tally & frequency · Languages — idioms and meaning

SIGNATURE MISSION T7-13 · The Unfair Machine

The class trains a deliberately one-sided suggestion machine — then watches unfairness walk in with the data. Bias stops being a vocabulary word and becomes something students caused, observed and can explain.

CH 3 AI at Work in India

T7-10 – T7-12

Where is AI already working — and how do I read its report card?

OUTCOMES

- Tour five sectors — healthcare, education, transport, agriculture, communication
- Write benefit-and-limit report cards for real AI helpers
- Draw bar, line and pie charts — and catch the charts that lie

CORE IDEAS

structured data · trend · sample · benefit · limit

Core understanding — Every real AI system deserves two grades: what it does well and where it stops — professionals always read both.

SUBJECTS

Maths — data handling & percentages · Social Science — community work

CH 4 When AI Gets It Wrong

T7-13 – T7-15

Who taught the machine — and who answers for its mistakes?

OUTCOMES

- Train a deliberately one-sided suggestion machine and watch bias walk in with the data
- Run three checks on faked or edited content, and hunt what is missing
- Debate who is responsible when a machine causes harm

CORE IDEAS

bias · misinformation · verify · fairness · accountability

Core understanding — A machine is a mirror of its training: unfair data makes unfair machines, and people — not machines — answer for the harm.

SUBJECTS

Social Science — fairness & rules · English — checking evidence

CH 5 Class Survey Machine CAPSTONE

T7-16 – T7-17

Can our class build — and check — its own prediction machine?

OUTCOMES

- Collect structured data with a real survey, chart it and read the pattern
- Predict a number from a trend, then check the sample for bias
- Grade the class machine against another section's test set

CORE IDEAS

structured data · sample · trend · prediction · report card

Core understanding — An honest prediction machine checks itself: sample, bias and test set are not extra work — they are the work.

SUBJECTS

Maths — data & trends · Social Science — surveys · English — reporting

BY THE END OF THIS YEAR

- Run classification, regression and clustering — and explain train-versus-test honestly
- Read any chart critically, and write a fair benefit-and-limit report card
- Trace an AI failure back to its data, and name who is responsible

1 BADGE EARNED · LEVEL 2 — SORT & PREDICT COMPLETE

CROSS-CURRICULAR FOOTPRINT

How many chapters reach each subject



CLASS 8

LEVEL 3 — BUILD THE CYCLE

Build the AI Cycle

17 PAPER MISSIONS · 5 CHAPTERS

My case file for planning, testing and improving an AI people can trust

YEAR 3 OF 5

Class 8 promotes the investigator to engineer — and engineers begin with plans, not machines. Students learn the loop every serious AI team runs: scope the problem, collect the right data, build and test, reflect and improve. They audit real tables for who is missing, repair unfair datasets, run a sampling game that ends with a surprise, hand-run the simplest classifier and grade its accuracy honestly, and put a text generator's fluency under the microscope to see why smooth talk can still invent facts. The capstone is a complete AI project proposal — problem, audited data plan, chosen tool, honest test, fairness check and known limits — pressure-tested by classmates in a live review gallery.

5	17	16	0
CHAPTERS	PAPER MISSIONS	WORDS TO KEEP	SCREENS NEEDED

CH 1 The AI Project Cycle

T8-01 – T8-03

What does it take to build an AI that people can trust?

OUTCOMES

- Turn vague complaints into precise, buildable problem statements
- Sequence the cycle — scope, data, build, reflect — and plan a dataset first
- Walk a plant-health case where skipping a step costs the crop

CORE IDEAS

machine learning · AI project cycle · problem statement · dataset

Core understanding — Serious AI is built as a loop, not a leap: scope the problem, plan the data, build, then reflect and improve.

SUBJECTS

Maths — percentages & data · English — precise definitions · Science — plant health case

CH 2 Data & Fairness

T8-04 – T8-07

If the data only knows some of us, who does the machine actually serve?

OUTCOMES

- Audit real tables — count who is over-represented and who is missing
- Run a sampling game that ends with a surprise, then repair an unfair dataset three ways
- Write the consent rules every honest data collector must follow

CORE IDEAS

sampling · consent · fairness · representation

Core understanding — A dataset is a promise about who the system will serve — engineers must check who the data includes before any model is built.

SUBJECTS

Maths — percentages, sampling · Social Science — inclusion & community · English — persuasive asking

SIGNATURE MISSION **T8-08 · Nearest Neighbour by Hand**

Students hand-run the simplest learning machine — a classifier that judges new cases by their most similar labelled examples — then grade it with a real accuracy score on a hidden test set, and face the limits it reveals.

CH 3 How Models Learn & What No-Code Tools Do

T8-08 – T8-10

How does a machine decide — and how do we measure whether it decides well?

OUTCOMES

- Hand-run a nearest-neighbour classifier that judges new cases by similar examples
- Grade it with a real accuracy percentage on a hidden test set
- Survey the no-code market — chatbots, image classifiers, predictors — with their limits

CORE IDEAS

nearest neighbour · accuracy · no-code tool · classifier

Core understanding — "The model works" means nothing until accuracy is measured on data the model has never seen — and even then, read the limits.

SUBJECTS

Maths — distance on a grid, percentages · Science — leaf features · Social Science — honest advertising

CH 4 Ethics & Responsible AI

T8-11 – T8-14

Why does a machine that sounds so right sometimes get it so wrong — and who answers?

OUTCOMES

- Play the next-word prediction game and see why fluency invites invention
- Hunt hallucinations in a real-looking AI paragraph
- Take an AI scenario to court as judge, lawyer and witness — then write a personal AI-use pledge

CORE IDEAS

text generator · hallucination · misinformation · responsible use

Core understanding — A text generator predicts, it does not know; the very trick that makes it fluent is the trick that makes it invent.

SUBJECTS

English — language patterns & argument · Social Science — privacy & public impact · Maths — likelihood

CH 5 Capstone: The AI Project Proposal CAPSTONE

T8-15 – T8-17

Can you plan an AI that people can trust — and defend the plan to a live audience?

OUTCOMES

- Write a complete proposal: problem, audited data plan, chosen tool, honest test, fairness check
- Survive the peer-review gallery as classmates pressure-test every claim
- Sharpen real budget choices with a trade-off and optimisation round

CORE IDEAS

peer review · trade-off · optimise · stakeholder

Core understanding — A proposal that cannot survive friendly criticism is not ready for the world — review is not a hurdle, it is the design tool.

SUBJECTS

All five strands by design · English — persuasive writing · Maths — budgeting & optimisation

BY THE END OF THIS YEAR

- Run the AI project cycle from sharp problem statement to fairness check
- Audit a dataset for representation and consent — and repair what is unfair
- Measure accuracy honestly, and name a tool's limits before praising it

1 BADGE EARNED · LEVEL 3 — BUILD THE CYCLE COMPLETE

CROSS-CURRICULAR FOOTPRINT

How many chapters reach each subject



CLASS 9

LEVEL 4 — REASON

The Logic Under the Magic

21 PAPER MISSIONS · 5 CHAPTERS

My case file for the maths that makes machines learn — and the reasons they fail

YEAR 4 OF 5

Class 9 opens the engine room. The analyst replaces admiration with mechanism: averages and the spread they hide, probability rolled on the desk, a line of best fit drawn through real dots, nearest-neighbour distances used as evidence. Data literacy turns students into inspectors of the numbers that feed learning systems — auditing sources, catching the three classic ways graphs lie, and separating privacy from security. Generative AI is built, not browsed: a next-word table constructed by hand turns the mystery into a mechanism before the ethics floor debates deepfakes, copyright and access. The capstone lifts eyes to the world: one SDG-linked AI brief, defended before an ethics review panel.

5	21	20	0
CHAPTERS	PAPER MISSIONS	WORDS TO KEEP	SCREENS NEEDED

CH 1 Reflect, Scope, Plan

T9-01 – T9-04

Why does AI work — and when does it fail?

OUTCOMES

- Separate rule-based systems from learning-based ones, and nest AI, ML and deep learning
- Rebuild the planning toolkit at analyst level: 4Ws canvas, stakeholders, system map
- Name every data feature a system needs before any mathematics begins

CORE IDEAS

deep learning · rule-based system · learning-based system · problem scoping · system map

Core understanding — Before the mathematics comes the map: a system that is scoped clearly is already halfway to working honestly.

SUBJECTS

Maths — data & coordinates · Social Science — stakeholders · English — precise definitions

CH 2 Data Literacy

T9-05 – T9-09

Can I trust this number — and can I trust the picture it comes in?

OUTCOMES

- Sort data by type and structure, then audit a source for quality and motive
- Catch the three classic ways graphs lie — truncated axes, cherry-picked windows, impossible pies
- Separate privacy from security, and build a personal data dashboard

CORE IDEAS

qualitative data · quantitative data · unstructured data · privacy · security

Core understanding — A number is a claim, and a graph is an argument — analysts check both before they believe either.

SUBJECTS

Maths — statistics & graphs · Social Science — sources & media · English — critical reading

SIGNATURE MISSION T9-16 · How a Text Generator Guesses

Students build a next-word table with their own pencil and generate text from it. The mystery of fluent AI becomes a mechanism anyone can explain — which is what makes its inventiveness, and its honesty risks, discussable.

CH 3 The Maths Behind AI (on paper)

T9-10 – T9-14

What does the machine actually compute — and when does the computation fool it?

OUTCOMES

- Compute mean, median and mode — and the spread each one hides
- Roll probability on the desk; draw and interrogate a line of best fit
- Classify with nearest-neighbour distances, then turn a table into an algorithm

CORE IDEAS

mean · median · mode · probability · correlation · line of best fit · nearest neighbour

Core understanding — Every learning machine computes arithmetic a student can do by hand — doing it once by hand is worth a hundred demos.

SUBJECTS

Maths — statistics, probability, coordinate geometry · Science — experiment & error · English — precise claims

CH 4 Generative AI & Ethics

T9-15 – T9-18

Machines that write, speak and show faces — how do they work, and what do we owe each other?

OUTCOMES

- Build a next-word table by hand and generate text from it — mystery becomes mechanism
- Run real-or-AI evidence checks; examine deepfakes, copyright and academic integrity
- Hold the balloon debate: which AI future deserves the seat?

CORE IDEAS

generative AI · deepfake · verification routine · bias

Core understanding — Generative AI does not retrieve answers — it constructs plausible ones; verifying what it makes is now a core life skill.

SUBJECTS

English — language patterns & argument · Social Science — access & equity · Art — what is authorship?

CH 5 Capstone: The SDG-linked AI Brief CAPSTONE

T9-19 – T9-21

Can you connect an AI idea to the world's goals — and defend it before a panel?

OUTCOMES

- Connect a scoped AI idea to one Sustainable Development Goal
- Write a formal brief with honest data and verified claims
- Face the ethics review panel — arguing from every stakeholder's chair

CORE IDEAS

SDG · ethics review · flowchart · trace

Core understanding — Good ideas earn their keep under examination: an AI brief is only as strong as the ethics review it survives.

SUBJECTS

Social Science — SDGs & development · Maths — the Chapter 3 mechanisms · English — the brief as argument

BY THE END OF THIS YEAR

- Compute and question the numbers underneath any AI claim
- Audit sources, catch lying graphs, and keep private data safe
- Explain generative AI as a mechanism — and use it honestly

1 BADGE EARNED · LEVEL 4 — REASON COMPLETE

CROSS-CURRICULAR FOOTPRINT

How many chapters reach each subject



CLASS 10

LEVEL 5 — EVALUATE & DESIGN

Decide, Evaluate, Design

21 PAPER MISSIONS · 5 CHAPTERS

My field book for judging AI systems, designing them responsibly — and defending the call

YEAR 5 OF 5

Class 10 completes the journey from noticing to designing. The designer draws the map of every model family met since Class 6, runs a real kernel across an image by hand, and builds text tables a machine could count. Then the single number that flatters every product — accuracy — is replaced by a discipline: confusion matrices, precision, recall and base rates, chosen to match the stakes. Governance follows the machinery: bioethics in medical triage, a full algorithm audit, privacy rights and responsibilities, and generative AI as the integrity question of this decade. The capstone is a complete responsible-design brief — with an ethics gate that can stop the project — defended before a board of judges.

5	21	24	0
CHAPTERS	PAPER MISSIONS	WORDS TO KEEP	SCREENS NEEDED

CH 1 Models Explained

T10-01 – T10-04

What exactly is the machine doing under there?

OUTCOMES

- Sort ten real products into model families on one map
- Meet the three learning strategies — supervised, unsupervised, reinforcement
- Run all four learning jobs on a kirana-store dataset

CORE IDEAS

model family · supervised learning · unsupervised learning · reinforcement learning · weight

Core understanding — Every AI product since Class 6 lives on one map of model families — knowing the family tells you what it can and cannot do.

SUBJECTS

Maths — functions & relations · Science — neuron models · Economics — market behaviour

CH 2 Machines That See and Read

T10-05 – T10-09

What does a machine actually receive — and what does it do with it?

OUTCOMES

- Run a real 3 × 3 kernel across a 6 × 6 image and watch an edge announce itself
- Stack filters into feature maps; normalise and clean text until a machine can count it
- Build bag-of-words and TF-IDF tables with a pencil

CORE IDEAS

pixel · grayscale · kernel · feature map · normalisation

Core understanding — Images are grids of numbers and text is countable patterns — the machine's way of seeing is learnable, and learning it removes the magic.

SUBJECTS

Maths — matrices & weighted sums · Science — optics & the eye · Languages — morphology, root words

SIGNATURE MISSION T10-13 · "99% Accurate"? The Base-Rate Trap

Students dismantle a "99% accurate" screening claim using the base rate, and watch an excellent-looking tool fail the very people it was built to help. Accuracy alone never again gets to end an argument.

CH 3 Judging a Model

T10-10 – T10-14

Is it any good — good for whom, at what, judged how?

OUTCOMES

- Split data for an honest test and count the four fates of every prediction
- Compute precision and recall — then choose the metric the stakes demand
- Dismantle a "99% accurate" claim with the base rate

CORE IDEAS

train/test split · confusion matrix · precision · recall · base rate

Core understanding — Accuracy alone is the least examined number in AI — professionals pick the metric that matches the stakes, and defend the choice.

SUBJECTS

Maths — percentages, ratios, conditional probability · Science — screening tests · Civics — accountability

CH 4 Ethics, Governance & Society

T10-15 – T10-18

Who decides — and who answers when the decision harms?

OUTCOMES

- Run a medical-triage tool through the four bioethics principles
- Conduct a full algorithm audit — bias, transparency, accountability — with evidence
- Map privacy rights against responsibilities; face generative AI as the integrity question

CORE IDEAS

accountability · transparency · data minimisation · integrity

Core understanding — Every AI decision has a person behind it — governance is how societies make sure that person can be named, questioned and answered.

SUBJECTS

Biology — medical ethics · Civics — rights & governance · Economics — labour markets · English — reasoned argument

CH 5 Capstone: The Responsible AI Design Brief CAPSTONE

T10-19 – T10-21

Can you design AI that survives its own audit?

OUTCOMES

- Design a complete system on paper — problem, consented data plan, model family, defensible metric
- Set deployment limits and an ethics gate that can stop the project
- Pitch it — naming its own doubts before the Board of Judges does

CORE IDEAS

design brief · ethics gate · deployment limits · pitch

Core understanding — Responsible design is not a section of the proposal — it is the proposal; a system that cannot survive its own audit is not finished.

SUBJECTS

All subjects by design · English — the pitch is reasoned writing · Economics — budgets and trade-offs

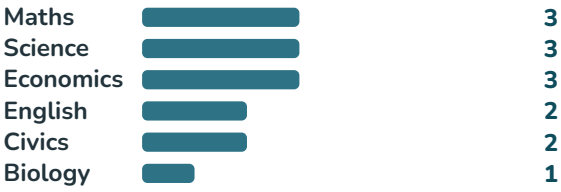
BY THE END OF THIS YEAR

- Judge any model with the metric the stakes demand — and defend the choice
- Audit a system for bias, transparency and accountability
- Design AI responsibly, and argue the call when the room pushes back

1 PORTFOLIO EARNED · LEVEL 5 — EVALUATE & DESIGN COMPLETE

CROSS-CURRICULAR FOOTPRINT

How many chapters reach each subject



The Five-Year Pathway at a Glance

Five books, one climb. Each year revisits the same five ideas — data, learning, fairness, evaluation and design — one level deeper, in a new role, with harder missions and more honest questions. The subjects already on the timetable carry the content: AI becomes a lens through which mathematics, language, science and social studies are used, not an extra subject to squeeze in.

LEVEL	CLASS	STUDENTS BECOME	YEAR FOCUS	CAPSTONE
1 · NOTICE	6	AI Detective	Spotting smart machines; data as the food of AI; patterns and predictions; safety and honesty online	My Neighbourhood AI Investigation
2 · SORT & PREDICT	7	AI Investigator	Running classification, regression and clustering; machines that see, read and recommend; bias, fakes and accountability	Class Survey Machine
3 · BUILD THE CYCLE	8	AI Engineer	The AI project cycle; data audits, sampling and consent; accuracy and no-code tools; responsible use	The AI Project Proposal
4 · REASON	9	AI Analyst	The mathematics under AI; data literacy and lying graphs; generative AI as mechanism and ethics	The SDG-linked AI Brief
5 · EVALUATE & DESIGN	10	AI Designer	Model families; kernels and text tables; precision, recall and base rates; governance and responsible design	The Responsible AI Design Brief

HOW THE SPIRAL WORKS

Each book stands alone, but the journey does not: Class 7 opens by using what Class 6 noticed, and Class 10 audits everything the four earlier years built. Skills return at greater depth — charts in Class 6 become confusion matrices in Class 10.

HOW LEARNING IS ASSESSED

Formative only. Every mission ends in reflection, every chapter in a recap, and every year in a badge — or, in Class 10, a portfolio. Capstones are judged live by classmates, panels and boards, so students practise defending evidence, not memorising answers.

WHAT A SCHOOL NEEDS

Printed handouts, pencils, and a curious classroom. No laboratory, no screens, no annual licence — every child receives the same book and the same chance. A year of AI literacy costs the price of a workbook, because opportunity should never depend on infrastructure.

Bring the AI Curriculum to your school

The five handouts in this series — one per class, 6 to 10 — are ready for adoption. Review a full copy, speak with our academic and psychology teams, and let us plan the rollout together.